

Cloud security

8

Security has been a concern since the early days of computing, when a computer was isolated and threats could only be posed by someone with access to the computer room. Once computers were able to communicate with one another, the Pandora's box of threats was wide open. In an interconnected world, various embodiments of malware can migrate easily from one system to another, cross national borders, and infect systems all over the world.

Security of computer and communication systems takes on a new urgency as the society becomes increasingly more dependent on the information infrastructure. Even the critical infrastructure of a nation can be attacked by exploiting flaws in computer security and often human naivety. Malware, such as the Stuxnet virus, targets industrial systems controlled by software [102]. The concept of *cyberwarfare* meaning “actions by a nation-state to penetrate another nation’s computers or networks for the purposes of causing damage or disruption” [109] was recently added to the dictionary.

A computer cloud is a target-rich environment for malicious individuals and criminal organizations. It is, thus, no surprise that security is a major concern for existing users and for potential new users of cloud computing services. Some of the security risks faced by computer clouds are shared with other systems supporting network-centric computing and network-centric content, e.g., service-oriented architectures, grids, and web-based services.

Cloud computing is an entirely new computing paradigm based on a new technology. It is therefore reasonable to expect that new methods to deal with some of existing security threats will be developed, while other perceived threats will prove to be exaggerated [25]. Indeed, “early on in the life cycle of a technology, there are many concerns about how this technology will be used...they represent a barrier to the acceptance...over the time, however, the concerns fade, especially if the value proposition is strong enough” [248].

The breathtaking pace of developments in information science and technology has many side effects. One of them is that standards, regulations, and laws governing the activities of organizations supporting the new computing services, and in particular utility computing, have yet to be adopted. As a result, many issues related to privacy, security, and trust in cloud computing are far from being settled. For example, there are no international regulations related to data security and privacy. Data stored on a computer cloud can freely cross national borders throughout the data centers of the CSP.

The chapter starts with a discussion of cloud users concerns related to security in Section 8.1. In Section 8.2, we elaborate on the security threats perceived by cloud users, as mentioned in Section 2.10. Privacy and trust are covered in Sections 8.4 and 8.5. Encryption discussed in Section 8.6 protects data in cloud storage, but data must be decrypted for processing. Threats during processing caused by flaws in the hypervisors, rogue VMs, or a VMBR, discussed in Section 5.15, cannot be ignored.

Analysis of database service security in Section 8.7 is followed by a presentation of operating system security, VM security, and the security of virtualization in Sections 8.8, 8.9, and 8.10, respectively.

Sections 8.11 and 8.12 analyze the security risks posed by shared images and by a management OS. An overview of the Xoar hypervisor, a version of Xen that breaks the monolithic Design of the TCB, is discussed in Section 8.13, followed by mobile device security in Section 8.14. Section 8.15 and 8.16 cover the current cloud security threats and AWS security.

8.1 Security—the top concern for cloud users

The idea that moving to a cloud liberates an organization from concerns related to computer security and eliminates a wide range of threats to data integrity is accepted by some members of the IT community. Some argue that cloud security is in the hands of experts, hence cloud users are even better protected than before. As we shall see throughout this chapter, these seem to be rather naive points of view. Outsourcing computing to a cloud generates major new security and privacy concerns. Moreover, the Service Level Agreements do not provide adequate legal protection for cloud computer users who are often left to deal with events beyond their control.

Previously, cloud users were accustomed to operate inside a secure perimeter protected by a corporate firewall. Now, they have to extend their trust to the cloud service provider if they wish to benefit from the economical advantages of utility computing. The transition from a model where users have full control of all systems and where their sensitive information is stored and processed is a difficult one. The reality is that virtually all surveys report that security is the top concern of cloud users.

The major user concerns are: unauthorized access to confidential information and data theft. Data is more vulnerable in storage than while it is being processed. Data is kept in storage for extended periods of time, while during processing, it is exposed to threats for a relatively short time. Hence, close attention should be paid to the security of storage servers and to data in transit. There is also the risk of unauthorized access and data theft posed by rogue employees of a CSP. Cloud users are concerned about insider attacks because hiring and security screening policies of a CSP are totally opaque processes.

The next concern regards user control over the lifecycle of data. It is virtually impossible for a user to determine if data that should have been deleted actually was. Even if deleted, there is no guarantee that the media was totally wiped out and that the next user is not able to recover confidential data. This problem is exacerbated because the CSPs rely on seamless backups to prevent accidental data loss. Such backups are done without user knowledge or consent. During this exercise data records can be lost, accidentally deleted, or accessible to an attacker.

Lack of standardization is next on the list of concerns. Today, there are no interoperability standards, as we have discussed in Section 2.6. Many questions do not have satisfactory answers at this time; for example: What can be done when the service provided by the CSP is interrupted? How should one access critically needed data in case of a blackout? What if the CSP drastically raises its prices? What is the cost of moving to a different CSP?

It is undeniable that auditing and compliance pose an entirely different set of challenges in cloud computing. These challenges are not yet resolved. A full audit trail on a cloud is an unfeasible proposition at this time.

Another, less analyzed, user concern is that cloud computing is based on a new technology expected to evolve into the future. Case in point, autonomic computing is likely to enter the scene. When this happens, self-organization, self-optimization, self-repair, and self-healing could generate additional se-

curity threats. In an autonomic system, it will be even more difficult than at the present time to determine when an action occurred, what was the reason for that action, and if and how it created the opportunity for an attack or for data loss. It is still unclear how autonomic computing can be compliant with privacy and legal issues.

There is no doubt that multitenancy is the root cause of many user concerns. Nevertheless, multitenancy enables a higher server utilization, thus lower costs. The users have to learn to live with multitenancy, one of the pillars of utility computing. The threats caused by *multitenancy* differ from one cloud delivery model to another. For example, in the case of SaaS, private information such as names, addresses, phone numbers, and possibly credit card numbers of many users are stored on one server; when the security of that server is compromised, a large number of users are affected.

Users are greatly concerned about the legal framework for enforcing cloud computing security. The cloud technology has moved much faster than cloud security, and privacy legislators and users have legitimate concerns regarding the ability to defend their rights. As the data centers of a CSP may be located in several countries, it is difficult to understand which laws apply: the laws of the country where information is stored and process; the laws of the countries the information crossed when sent by the user; or the laws of the user's country.

To further complicate legal aspects of cloud computing security, a CSP may outsource handling of personal and/or sensitive information. Existing laws stating that the CSP must exercise reasonable security may be difficult to implement when there is a chain of outsourcing to companies in different countries. Lastly, a CSP may be required to share with law enforcement agencies private data. For example, Microsoft was served a subpoena to provide emails exchanged by users of the Hotmail service.

The question is: What can and should cloud users do to minimize the security risks regarding the data handling by the CSP? First, a user should evaluate the security policies and the mechanisms the CSP has in place to enforce these policies. Then, the user should analyze the information that would be stored and processed on the cloud. Finally, contractual obligations should be clearly spelled out. The contract between the user and the CSP should [396]:

1. State explicitly CSP's obligations for handling securely sensitive information and to comply with privacy laws.
2. Spell out CSP liabilities for mishandling sensitive information, e.g., data loss.
3. Spell out the rules governing ownership of the data.
4. Specify the geographical regions where information and backups can be stored.

To minimize security risks, a user may try to avoid processing sensitive data on a cloud. Google's Secure Data Connector carries out an analysis of the data structures involved and allows access to data protected by a firewall. This solution is not feasible for some applications, e.g., processing of medical or personnel records; it may not be feasible when the cloud processing workflow requires cloud access to the entire volume of user data. When the volume of sensitive data or the processing workflow requires sensitive data to be stored on the cloud then, whenever feasible, data should be encrypted [192,526].

8.2 Cloud security risks

Some believe that it is very easy, possibly too easy, to start using cloud services without a proper understanding of the security risks and without the commitment to follow user responsibilities spelled

out by Service Level Agreements (SLA).¹ A first question is: What are the security risks faced by cloud users? A cloud could be used to launch large-scale attacks against other components of the cyber infrastructure, so the next question is: How can nefarious use of cloud resources be prevented?

There are multiple ways to look at the security risks for cloud computing. Three broad classes of security risks for cloud computing are: traditional security threats, threats related to system availability, and threats related to third-party data control [107].

Traditional threats are those that have been experienced for some time by any system connected to the Internet, but with some cloud-specific twists. The impact of traditional threats is amplified due to the vast amount of cloud resources and the large user population that can be affected. The long list of cloud user concerns includes also the fuzzy bounds of responsibility between the providers of cloud services and users and the difficulties to accurately identify the cause of a problem. The traditional threats begin at the user site. The user must protect the infrastructure used to connect to the cloud and to interact with the application running on the cloud. This task is more difficult because some components of this infrastructure are outside the firewall protecting the user.

The next threat is related to the authentication and authorization process. The procedures in place for one individual do not extend to an enterprise. In this case, the cloud access of the members of an organization must be nuanced; various individuals should be assigned distinct levels of privilege based on their role in the organization. It is also nontrivial to merge or adapt the internal policies and security metrics of an organization with the ones of the cloud.

Moving from the user to the cloud, we see that the traditional attacks have already affected cloud service providers. The favorite means of attack are: distributed denial of service (DDoS) attacks that prevent legitimate users to access cloud services, phishing, SQL injection, or cross-site scripting. Phishing is an attack intended to gain information from a database by masquerading as a trustworthy entity. Such information could be names and credit card numbers, social security numbers, and other personal information stored by online merchants or other service providers.

SQL injection is a form of attack typically used against a web site. In this case, an SQL command entered in a web form causes the contents of a database used by the web site to be either dumped to the attacker or altered. SQL injection can be used against other transaction processing systems, and it is successful when the user input is not strongly typed or rigorously filtered. Cross-site scripting is the most popular form of attack against web sites; a browser permits the attacker to insert client-scripts into the web pages, and thus bypass the access controls at the web site.

Identifying the path followed by an attacker is much more difficult in a cloud environment. Cloud servers host multiple VMs, and multiple applications may run under one VM. Multitenancy in conjunction with hypervisor vulnerabilities could open new attack channels for malicious users. Traditional investigation methods based on digital forensics cannot be extended to a cloud where the resources are shared among a large user population and the trace of events related to a security incident is wiped out due to the high rate of write operations on any storage media.

Availability of cloud services is another concern. System failures, power outages, and other catastrophic events could shutdown cloud services for extended periods of time. Data lock-in discussed in Section 2.6 could prevent a large organization whose business model depends on these data to function properly, when such a rare event occurs.

¹ AWS SLAs can be found at <https://aws.amazon.com/legal/service-level-agreements/>.

Clouds can also be affected by phase transition phenomena and other effects specific to complex systems. Another critical aspect of availability is that the users cannot be assured that an application hosted on the cloud returns correct results.

Third-party control generates a spectrum of concerns caused by a lack of transparency and limited user control. For example, a cloud provider may subcontract some resources from a third party whose level of trust is questionable. There are examples when subcontractors failed to maintain the customer data. There are also examples when the third party was not a subcontractor but a hardware supplier, and the loss of data was caused by poor-quality storage devices [107].

Storing proprietary data on the cloud is risky because cloud provider espionage poses real dangers. The terms of contractual obligations usually place all responsibilities for data security with the user. The Amazon Web Services customer agreement does not help user's confidence because it states "We...will not be liable to you for any direct, indirect, incidental...damages...nor...be responsible for any compensation, reimbursement, arising in connection with: (A) your inability to use the services...(B) the cost of procurement of substitute goods or services...or (D) any unauthorized access to, alteration of, or deletion, destruction, damage, loss or failure to store any of your content or other data."

It is very difficult for a cloud user to prove that data has been deleted by the service provider. The lack of transparency makes auditability a very difficult proposition for cloud computing. Auditing guidelines elaborated by the National Institute of Standards (NIST), such as the Federal Information Processing Standard (FIPS) and the Federal Information Security Management Act (FISMA), are mandatory for US Government agencies.

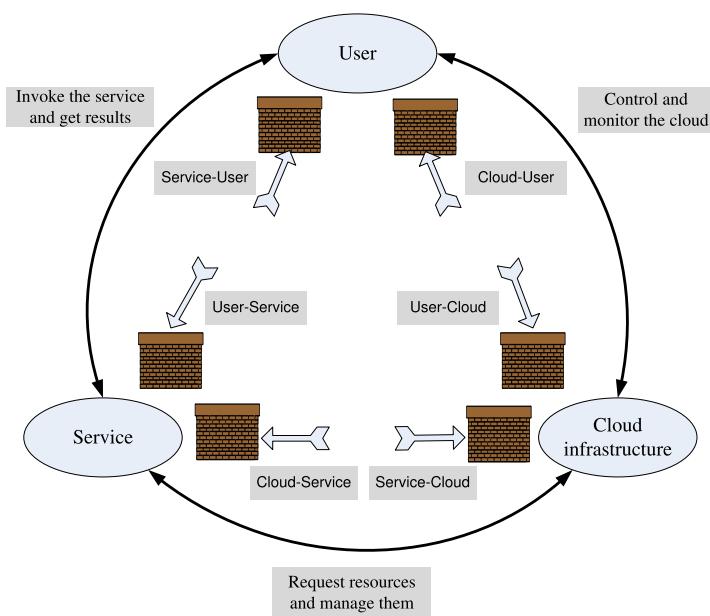
The 2010 Cloud Security Alliance (CSA) report. The report identifies seven top threats to cloud computing. These threats are: the abusive use of the cloud; APIs that are not fully secure, malicious insiders; shared technology; account hijacking; data loss or leakage; and unknown risk profile [121]. According to this report, the IaaS delivery model can be affected by all threats and PaaS can be affected by all, except the shared technology, while SaaS is affected by all, except abuse and shared technology.

Abusing the cloud refers to conducting nefarious activities from the cloud, for example, using multiple AWS instances or applications supported by IaaS to launch distributed denial of service attacks or to distribute spam and malware. *Shared technology* considers threats due to multitenant access supported by virtualization. Hypervisors can have flaws allowing a guest OS to affect the security of the platform shared with other VMs.

Insecure APIs may not protect the users during a range of activities, starting with authentication, and access control to monitoring and control of the application during runtime. The cloud service providers do not disclose their hiring standards and policies, thus the risks of *malicious insiders* cannot be ignored. The potential harm due to this particular form of attack is high.

Data loss and data leakage are two risks with devastating consequences for an individual or an organization using cloud services. Maintaining copies of the data outside the cloud is often unfeasible due to the sheer volume of data. If the only copy of the data is stored on the cloud, then sensitive data is permanently lost when cloud data replication fails, followed by a storage media failure. As some of the data often includes proprietary or sensitive data, access to such information by third parties could have severe consequences.

Account or service hijacking is a significant threat, and cloud users must be aware of and guard against all methods to steal credentials. Lastly, *unknown risk profile* refers to exposure to the ignorance or underestimation of the risks of cloud computing.

**FIGURE 8.1**

Surfaces of attacks in a cloud computing environment.

The 2011 CSA report. The report “Security Guidance for Critical Area of Focus in Cloud Computing V3.0” provides a comprehensive analysis of the risks and makes recommendations to minimize the risk in cloud computing [122].

An attempt to identify and classify the attacks in a cloud computing environment is discussed in [212]. The three actors involved in the model are: the user, the service, and the cloud infrastructure, and there are six types of attacks possible; see Fig. 8.1. The user can be attacked from two directions, the service and the cloud. Secure Sockets Layer (SSL) certificate spoofing, attacks on browser caches, or phishing attacks are example of attacks that originate at the service. The user can also be a victim of attacks that either truly originate or that spoof origination from the cloud infrastructure.

The service can be attacked by the user. Buffer overflow, SQL injection, and privilege escalation are the common types of attacks from the service. The service can also be the subject of attacks by the cloud infrastructure, and this is probably the most serious line of attack. Limiting access to resources, privilege-related attacks, data distortion, and injecting additional operations are only a few of the many possible lines of attacks originated at the cloud.

The cloud infrastructure can be attacked by a user that targets the cloud control system. These types of attacks are the same a user would direct toward any other cloud service. The cloud infrastructure may also be targeted by a service requesting an excessive amount of resources and causing the exhaustion of the resources.

Top cloud security threats. The 2016 CSA report lists top security threats [407]:

1. Data breaches. The most damaging breaches concern sensitive data, including financial and health information, trade secrets, and intellectual property. The ultimate responsibility rests with the organizations maintaining data on the cloud, and CSA recommends that organizations use multifactor authentication and encryption to protect against data breaches. Multifactor authentication, such as one-time passwords, phone-based authentication, and smart card protection, make it harder for attackers to use stolen credentials.
2. Compromised credentials and broken authentication. Such attacks are due to lax authentication, weak passwords, and poor key and/or certificate management.
3. Hacked interfaces and APIs. Cloud security and service availability can be compromised by a weak API. When third parties rely on APIs, more services and credentials are exposed.
4. Exploited system vulnerabilities. Resource sharing and multitenancy create new attack surfaces, but the cost to discover and repair vulnerabilities is small compared to the potential damage.
5. Account hijacking. All accounts should be monitored so that every transaction can be traced to the individual requesting it.
6. Malicious insiders. This threat can be difficult to detect, and system administrator errors could sometimes be falsely diagnosed as threats. A good policy is to segregate duties and enforce activities, such as logging, monitoring, and auditing administrator activities.

The other six threats are: advanced persistent threats (APTs), permanent data loss, inadequate diligence, cloud service abuse, DoS attacks, and shared technology.

An update of the “Cloud Controls Matrix” spells out the control specification impact on the cloud architecture, cloud delivery models, and other aspects of the cloud ecosystem; see <https://cloudsecurityalliance.org/download/cloud-controls-matrix-v3-0-1/>.

Cloud vulnerability incidents reported over a period of four years and data breach incidents in 2014 identified several other threats including [471]: hardware failures, natural disasters, cloud-related malware, inadequate infrastructure design and planning, point-of-sale (POS) intrusions and payment card skimmers, crimeware and cyber espionage, insider and privilege misuse, web app attacks, and physical theft/loss.

2021-global CSIO report. The report, commissioned by Dynatrace in 2021, stresses that the need to innovate faster and shift to cloud-native application architectures built on microservices, containers, and Kubernetes is driving complexity and creates vulnerability blind spots. It is becoming nearly impossible to distinguish between potential vulnerabilities and critical exposures. It is also virtually impossible for developers to conduct manual security scans or have the means to automatically identify potential vulnerabilities. Microservices are constantly changing and sometimes migrate from one cloud to another.

The report <https://assets.dynatrace.com/en/docs/report/2021-global-ciso-report.pdf> also acknowledges that “alternative approaches, such as scanning source code or container images in pre-production environments, cannot offer real-time visibility into live exposures. This means vulnerabilities often remain undetected in production, where attackers can exploit them.” The only hope for vulnerability management is future fully automated runtime security.

8.3 Security as a service (SecaaS)

SecaaS includes security services offered by a CSP. Several potential benefits of SecaaS [123] are: (i) intelligence-sharing since it protects multiple clients simultaneously and shares data intelligence and data across them; (ii) insulation of clients as it can intercept attacks before they hit an organization; and last but not least (iii) scaling and cost.

Potential concerns of SecaaS are: (i) regulation differences because the service may be unable to assure compliance in all jurisdictions an organization operates in; and (ii) lack of visibility as the service provider may not reveal details of how it implements its own security and manages its own environment. SecaaS offerings include:

1. *Identity, Entitlement, and Access Management Services* based on Federated Identity Brokers, strong authentication, multifactor authentication, etc.
2. *Cloud Access and Security Brokers*. Cloud Security Gateways connect to on-premises tools to help an organization detect, assess, and potentially block cloud usage and unapproved services.
3. *Web Security Gateways* providing an added layer of protection in addition to anti-malware software to prevent malware from entering the enterprise via activities such as web browsing.
4. *Email Security* including protection against phishing and malicious attachments, as well as enforcing corporate policies like acceptable use and spam prevention.
5. *Security Assessment* including traditional security/vulnerability assessments of assets, application security assessments, and cloud platform assessment tools.
6. *Web Application Firewalls*.
7. *Intrusion Detection/Prevention*.
8. *Security Information and Event Management*. Aggregates log and event data from virtual and real networks, applications, and systems and provides real-time reporting and alerts.
9. *Encryption and Key Management*.
10. *Business Continuity and Disaster Recovery*.
11. *Distributed Denial of Service Protection*.

CSA recommends that, before engaging a SecaaS, an organization should ensure that the service is compatible with the organization's own current and future plans, such as supported cloud platforms, the workstation, and the mobile operating systems it accommodates. Also, it is important to understand any security-specific requirements for data-handling and investigative and compliance support.

8.4 Privacy and privacy impact assessment

The term *privacy* refers to the right of an individual, a group of individuals, or an organization to keep information of a personal nature or proprietary information from being disclosed. Many nations view privacy as a basic human right. The Universal Declaration of Human Rights, article 12, states: "No one shall be subjected to arbitrary interference with his privacy, family, home or correspondence, nor to attacks upon his honor and reputation. Everyone has the right to the protection of the law against such interference or attacks."

U.S. Constitution contains no express right to privacy, however the Bill of Rights reflects framers concern for protecting specific aspects of privacy.² In the UK, privacy is guaranteed by the Data Protection Act. The European Court of Human Rights has developed several documents defining the right to privacy. Privacy laws differ from country to country; laws in one country may require public disclosure of information considered private in other countries and cultures.

At the same time, the right to privacy is limited by laws. For example, taxation laws require individuals to share information about personal income or earnings. Individual privacy may conflict with other basic human rights, e.g., with freedom of speech. The digital age has confronted legislators with significant challenges related to privacy as new threats have emerged. For example, personal information voluntarily shared, but stolen from sites granted access to it, or misused can lead to *identity theft*.

Some countries have been more aggressive in addressing the new privacy concerns than others. For example, European Union (EU) countries have strict laws governing handling of personal data in the digital age. A sweeping new privacy right, the “right to be forgotten,” is codified as part of a broad new proposed data-protection regulation in EU. This right addresses the following problem: It is very hard to escape your past now when every photo, status update, and tweet lives forever on some web site.

Our discussion targets primarily public clouds where privacy has an entirely new dimension because data, often in an unencrypted form, resides on servers owned by a CSP. Services based on individual preferences, location of individuals, membership in social networks, or other personal information present a special risk. The data owner cannot rely exclusively on the CSP to guarantee data privacy.

Privacy concerns are different for the three cloud delivery models and also depend on the actual context. For example, consider Gmail, a widely used SaaS delivery model; Gmail privacy policy reads (see <http://www.google.com/policies/privacy/> accessed on October 6, 2012): “We collect information in two ways: information you give us... like your name, email address, telephone number or credit card; information we get from your use of our services such as: ...device information, ...log information, ...location information, ...unique application numbers, ...local storage, ...cookies and anonymous identifiers... We will share personal information with companies, organizations or individuals outside of Google if we have a good-faith belief that access, use, preservation or disclosure of the information is reasonably necessary to: meet any applicable law, regulation, legal process or enforceable governmental request; ...protect against harm to the rights, property or safety of Google, our users or the public as required or permitted by law. We may share aggregated, non-personally identifiable information publicly and with our partners—like publishers, advertisers or connected sites. For example, we may share information publicly to show trends about the general use of our services.”

The main aspects of cloud privacy are: the lack of user control, potential unauthorized secondary use, data proliferation, and dynamic provisioning [396]. The lack of user control refers to the fact that user-centric data control is incompatible with cloud usage. Once data is stored on the servers of the CSP, the user loses control of the exact location and, in some instances, it could lose access to the data. For example, in the case of the Gmail service, the account owner has no control on where the data is stored or how long old emails are stored on some backups of the servers.

² The 1st Amendment covers the protection of beliefs, the 3rd Amendment privacy of homes, the 4th Amendment the privacy of persons and possessions against unreasonable searches, and the 5th Amendment the protection from self-incrimination. Thus, the privacy of personal information, and, according to some justices, the 9th Amendment that reads “The enumeration in the Constitution, of certain rights, shall not be construed to deny or disparage others retained by the people” can be viewed as protection of privacy in ways not explicitly specified by the first eight amendments in the Bill of Rights.

A CSP may obtain revenues from unauthorized secondary usage of the information, e.g., for targeted advertising. There are no technological means to prevent this use. Dynamic provisioning refers to threats due to outsourcing. A range of issues are very fuzzy, e.g.: how to identify the sub-contractors of a CSP, what rights to the data they have, and what rights to data are transferable in case of a bankruptcy or merger.

It is imperative for legislative bodies to address the multiple aspects of privacy in the digital age. A document elaborated by the Federal Trading Commission for the US Congress states [172]: “Consumer-oriented commercial web sites that collect personal identifying information from or about consumers online would be required to comply with the four widely-accepted fair information practices:

1. Notice—web sites should be required to provide consumers clear and conspicuous notice of their information practices, including what information they collect, how they collect it (e.g., directly or through nonobvious means such as cookies), how they use it, how they provide Choice, Access, and Security to consumers, whether they disclose the information collected to other entities, and whether other entities are collecting information through the site.
2. Choice—web sites should be required to offer consumers choices as to how their personal identifying information is used beyond the use for which the information was provided, e.g., to consummate a transaction. Such choices would encompass both internal secondary uses (such as marketing back to consumers) and external secondary uses, such as disclosing data to other entities.
3. Access—web sites would be required to offer consumers reasonable access to the information a web site has collected about them, including a reasonable opportunity to review information and to correct inaccuracies or delete information.
4. Security—web sites would be required to take reasonable steps to protect the security of the information they collect from consumers. The Commission recognizes that the implementation of these practices may vary with the nature of the information collected and the uses to which it is put, as well as with technological developments. For this reason, the Commission recommends that any legislation be phrased in general terms and be technologically neutral. Thus, the definitions of fair information practices set forth in the statute should be broad enough to provide flexibility to the implementing agency in promulgating its rules or regulations.”

There is the need for tools capable of identifying privacy issues in information systems, the so-called *Privacy Impact Assessment (PIA)*. As of mid-2017, there are no international standards for such a process, though different countries and organization require PIA reports. An example of an analysis is to assess the legal implications of the UK–US Safe Harbor process to allow US companies to comply with the European Directive 95/46/EC³ on the protection of personal data.

Such an assessment forces a proactive attitude towards privacy. An ab initio approach for embedding privacy rules in new systems is preferable to painful changes that could affect the functionality of existing systems. A PIA tool that could be deployed as web-based service proposed in [469] includes: project information, an outline of project documents, privacy risks, and stakeholders. The tool will produce a PIA report consisting of a summary of findings, a risk summary, security, transparency, and cross-borders data flows.

³ See <http://eur-lex.europa.eu/LexUriServ/LexUriServ.do?uri=CELEX:31995L0046:en:HTML>.

The centerpiece of the PIA tool is a knowledge base (KB) created and maintained by domain experts. The users of the SaaS service providing access to the PIA tool must complete a questionnaire. The system uses templates to generate additional questions necessary to complete the PIA report. An expert system infers which rules are satisfied by the facts in the database as provided by the users and executes the rule with the highest priority.

8.5 Trust

In the context of cloud computing trust is intimately related to the general problem of trust in online activities. We first discuss the traditional concept of trust and then the trust in online activities.

Trust. According to the Merriam–Webster dictionary, trust means “assured reliance on the character, ability, strength, or truth of someone or something.” Trust is a complex phenomenon; it enables cooperative behavior, promotes adaptive organizational forms, reduces harmful conflict, decreases transaction costs, facilitates formulation of ad hoc work groups, and promotes effective responses to crisis [424].

Two conditions must exist for trust to develop. The first is *risk*, the perceived probability of loss. Indeed, trust would not be necessary if there is no risk involved, if there is a certainty that an action will succeed. The second is *interdependence*, the interests of one entity cannot be archived without reliance on other entities. A trust relationship goes through three phases: (i) building phase, when trust is formed; (ii) stability phase, when trust exists; and (iii) dissolution phase, when trust declines.

There are various reasons for and forms of trust. Utilitarian reasons could be based on the belief that the costly penalties for breach of trust exceed any potential benefits from opportunistic behavior. This is the essence of *deterrence-based* trust. Another reason is the belief that the action involving the other party is in the self-interest of that party. This is the so-called *calculus-based* trust. After a long sequence of interactions *relational trust* between entities can be developed based on the accumulated experience of dependability and reliance on each other.

The common wisdom is that an entity must work very hard to build trust, but may lose the trust very easily. A single violation of trust can lead to irreparable damage. *Persistent trust* is based on the long-term behavior of an entity, while *dynamic trust* is based on a specific context, e.g., the state of the system or the effect of technological developments.

Trust in the Internet The trust in the Internet “obscures or lacks entirely the dimensions of character and personality, nature of relationship, and institutional character” of the traditional trust [365]. The missing identity, personal characteristics, and role definitions are elements we have to deal in the context of online trust.

The Internet offers individuals the ability to obscure or conceal their identities. The resulting anonymity reduces the clues normally used in judgments of trust. Identity is critical for developing trust relationships; it allows us to base our trust on the past history of interactions with an entity. Anonymity causes mistrust because identity is associated with accountability, and in the absence of identity, accountability cannot be enforced.

The opacity extends immediately from identity to personal characteristics. It is impossible to infer if the entity or individual we transact with is who it pretends to be because the transactions occur between entities separated in time and distance. Lastly, there are no guarantees that the entities we transact with fully understand the role they have assumed.

To remedy the loss of clues, we need security mechanisms for access control, transparency of identity, and surveillance. The mechanisms for access control are designed to keep out intruders and mischievous agents. Identity transparency requires that the relationship between a virtual agent and a physical person should be carefully checked security through methods such as biometric identification. Digital signatures and digital certificates are used for identification. Surveillance could be based on *intrusion detection* or can be based on logging and auditing. The first is based on real-time monitoring, while the second relies on offline sifting through audit records.

Credentials are used when the entity is not known; credentials are issued by a trusted authority and describe the qualities of the entity using the credential. A DDS (Doctor of Dental Surgery) diploma hanging on the wall of a dentist's office is a credential that the individual has been trained by an accredited university and hence capable to perform a set of procedures. A digital signature is a credential used in many distributed applications.

Policies and *reputation* are two ways of determining trust. Policies reveal the conditions to obtain trust and the actions when some of the conditions are met. Policies require the verification of credentials. Reputation is a quality attributed to an entity based on a relatively long history of interactions or possibly observations of the entity. Recommendations are based on trust decisions made by others and filtered through the perspective of the entity assessing the trust.

In a computer science context, “trust of a party A to a party B for a service X is the measurable belief of A in that B behaves dependably for a specified period within a specified context (in relation to service X),” [374]. An assurance about the operation of particular hardware or software component leads to persistent social-based trust in that component.

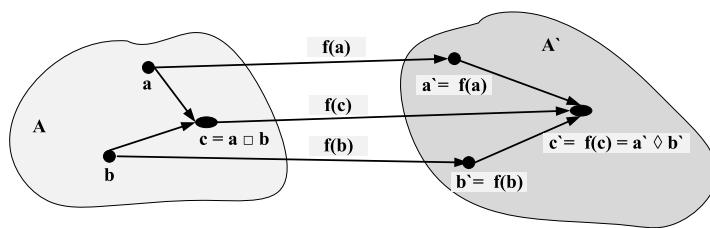
A comprehensive discussion of trust in computer services in the semantic web can be found in [33]. In Section A.1, we discuss the concept of trust in the context of cognitive radio networks where multiple transmitters compete for communication channels. Then, in Section A.2, we present a cloud-based trust management service.

8.6 Cloud data encryption

Individual users, large corporations, and the government ponder if it is safe to store sensitive information on a public cloud. Encryption is the obvious solution to protect outsourced data. The seminal RSA paper [418] and the survey of existing public-key cryptosystems in [429] are some of the notable publications in the vast literature dedicated to cryptosystems.

Cloud service providers offer encryption services. Amazon offers AWS Key Management Service (KMS) to create and control the encryption keys used by clients to encrypt their data. KMS is integrated with other AWS services including EBS, S3, RDS, Redshift, Elastic Transcoder, and WorkMail. AWS also offers Encryption SDK for developers.

Several recent research results in cryptography are important to data security in cloud computing. In 1999, Pascal Paillier proposed a trapdoor mechanism for public-key cryptosystems based on composite degree residuosity for hard-to-factor numbers $n = pq$ where p and q are large prime numbers [386]. This solution exploits the homomorphic properties of composite residuosity classes to design distributed cryptographic protocols. A major breakthrough are the algorithms for Fully Homomorphic Encryption (FHE) proposed by Craig Gentry in his seminal 2009 dissertation at Stanford University

**FIGURE 8.2**

A homomorphism $f : A \rightarrow A'$ is a structure-preserving map between sets A and A' with the composition operations \square and \diamond , respectively. Let $a, b, c \in A$ with $c = a \square b$ and $a', b', c' \in A'$ with $c' = a' \diamond b'$. Let $a' = f(a)$, $b' = f(b)$, $c' = f(c)$ be the results of the mapping $f(\cdot)$. The operation \diamond in the target domain produces the same result as the mapping the result of the operation \square applied to the two elements in the original domain, $f(a) \diamond f(b) = f(a \square b)$.

[192,193]. In recent years, searchable symmetric encryption protocols have been reported in [85] and [168].

Sensitive data is safe while in storage, provided that it is encrypted with strong encryption. But encrypted data must be decrypted for processing, and this opens a window of vulnerability. So, here is the first question examined in this section: Is it feasible to operate on encrypted data? The *homomorphic encryption*, a long-time dream of security experts, reflects the concept of homomorphism, a structure-preserving map $f(\cdot)$ between two algebraic structures of the same type; see Fig. 8.2.

When $f(\cdot)$ in Fig. 8.2 is a one-to-one mapping, the inverse of $f(\cdot)$ is $f^{-1} : A' \rightarrow A$ and $a = f^{-1}(a')$, $b = f^{-1}(b')$, $c = f^{-1}(c')$. In this case, we can carry out the operation \diamond in the target domain and apply the inverse mapping to get the same result produced by the \square operation in the original domain, $f^{-1}(a) \diamond f^{-1}(b) = f(a \square b)$. In the case of homomorphic encryption, the map is a one-to-one transformation, $f(\cdot)$ is the encryption procedure, its inverse, $f^{-1}(\cdot)$, is the decryption procedure, and the composition operation can be any arithmetic and logic operation. The homomorphism shows that we can conduct arithmetic and/or logic operations with encrypted data and the decryption of the result of these operation replicates the result of carrying out the same operations with the plaintext data. This implies that the window of vulnerability created when data is decrypted for processing disappears.

General computations with encrypted data are theoretically feasible using FHE algorithms. Unfortunately, the homomorphic encryption is not a practical solution at this time. Existing algorithms for homomorphic encryption increase the processing time with encrypted data by many orders of magnitude compared with the processing of plaintext data. A recent implementation of FHE [223] requires about six minutes per batch; the processing time for a simple operation on encrypted data dropped to almost one second after improvements in other experiments [153].

Users send a variety of queries to many large databases stored on clouds. Such queries often involve logic and arithmetic functions, so an important question is if it is feasible and practical to search encrypted databases. Application of widely used encryption techniques to database systems could lead to significant performance degradation. For example, if an entire column of a NoSQL database table contains sensitive information and is encrypted, then a query predicate with a comparison operator requires a scan of the entire table to evaluate the query. This is due to the fact that existing encryption algorithms do not preserve order, and database indices such as B-tree can no longer be used.

Order Preserving Encryption (OPE) [66] can be used for encryption of numeric data. OPE maps a range of numerical values into a much larger and sparse range of values. Let a order-preserving function $f : \{1, \dots, M\} \rightarrow \{1, \dots, N\}$ with $N \gg M$ be uniquely represented by a combination of M out of N ordered items. Given N balls in a bin, M black and $N - M$ white, we draw a ball at random without replacement at each step. The random variable X describing the total number of balls in our sample after we collect the k -th black ball follows the negative hypergeometric distribution. One can show that an order preserving $f(x)$ for a given point $x \in \{1, \dots, M\}$ has an NHG distribution over a random choice of f .

To encrypt plaintext x , the OPE encryption algorithm performs a binary search down to x . Given the secret key K , the algorithm first assigns $\text{Encrypt}(K, M/2)$, then $\text{Encrypt}(K, M/4)$ if the index $m < M/2$ and $\text{Encrypt}(K, 3M/4)$ otherwise, and so on, until $\text{Encrypt}(K, x)$ is assigned. Each ciphertext assignment is made according to the output of the negative hypergeometric sampling algorithm. One can prove by strong induction on the size of the plaintext space that the resulting scheme induces a random order-preserving function from the plaintext to ciphertext space.

It is sufficient to have an order-preserving hash function H (not necessarily invertible) to allow efficient range queries on encrypted data. The algorithm would use a secret key $(K_{\text{Encrypt}}, K_H)$, where K_{Encrypt} is a key for a normal (randomized) encryption scheme and K_H is a key for H . Then $\text{Encrypt}(K_{\text{Encrypt}}, x) \parallel H(K_H, x)$ will be the encryption of x [66].

Searching encrypted databases is of particular interest [13]. Several types of searches are frequently conducted including: single-keyword, multikeyword, fuzzy-keyword, ranked, authorized, and verifiable search. Searchable symmetric encryption (SSE) is used when an encrypted databases \mathcal{E} is outsourced to a cloud or to a different organization. SSE hides information about the database and the queries.

A client only stores the cryptographic key. To search a database, the client encrypts the query, sends it to the database server, receives the encrypted result of the query, and decrypts it using the cryptographic key. Information leakage from these searches is confined to query patterns, while disclosure of explicit data and query plaintext values is prevented.

An SSE protocol supporting conjunctive search and general Boolean queries on symmetrically encrypted data was proposed in [85]. This SSE protocol scales to very large databases. It can be used for arbitrarily structured data including free text search with the moderate and well-defined leakage to the outsourced server. The performance results of a prototype applied to encrypted search over the entire English Wikipedia are reported. The protocol was extended with support for range, substring, wildcard, and phrase queries [168].

The next question is if sensitive data stored on the servers of a private cloud is vulnerable. The threat posed by an outsider attacker is diminished if the private cloud is protected by an effective firewall. Nevertheless, there are dangers posed by an insider. If such an attacker has access to log files, it can infer the location of database hot spots, copy data selectively, and use the data for a nefarious activity. To minimize the risks posed by an insider, a set of protections rings should be enforced to restrict the access of each member of the staff to a limited area of the data base.

8.7 Security of database services

Cloud users who delegate the control of their data to the database services supported by virtually all CSP are concerned with Data Base as a Service (DBaaS) security aspects. The model used to evaluate

DBaaS security includes several group of entities: the data owners, the users of data, the cloud service providers, and third-party agents or Third Party Auditors (TPA).

Data owners and DBaaS users fear compromised integrity and confidentiality, as well as data unavailability. Insufficient authorization, authentication and accounting mechanisms, inconsistent use of encryption keys and techniques, alteration or deletion of records without maintaining backup, and operational failures are the major causes of data loss in DBaaS.

Some data integrity and privacy issues are due to the absence of authentication, authorization and accounting controls, or poor key management for encryption and decryption. Confidentiality means that only authorized users should have access to the data. Unencrypted data is vulnerable to bugs, errors, and attacks from external entities affecting data confidentiality.

Insider attacks are a concern for DBaaS users and data owners. Superusers have unlimited privileges, and misuse of superuser privileges poses a considerable threat to confidential data, such as medical records, sensitive business data, proprietary product data, etc. Malicious external attackers use spoofing, sniffing, man-in-the-middle attacks, side channeling, and illegal transactions to launch DoS attacks.

Another concern is illegal data recovery from storage devices, a side effect of multi-tenancy. CSP often carry out sanitation operations after deleting data from physical devices, but, unless a thorough scrubbing operation is carried out, sophisticated attackers can still recover information from storage devices. Data is also vulnerable during transfer from the data owner to the DBaaS through public networks. Encryption before data transmission can reduce the risks posed to the data in transit to the cloud.

Data provenance, the process of establishing the origin of data and its movement between databases, uses metadata to determine the data accuracy, but the security assessments are time-sensitive. Moreover, analyzing large provenance metadata graphs is computationally expensive.

Cloud users are not aware of the physical location of their data. This lack of transparency allows cloud service providers to optimize the use of resources, but in the case of security breaches, it is next to impossible for users to identify compromised resources. DBaaS users do not have fine-grained control of the remote execution environment and cannot inspect the execution traces to detect the occurrence of illegal operations.

To increase availability, performance, and to enhance reliability, DBaaS replicate data. Ensuring consistency among the replicas is challenging. Another critical function of DBaaS is to carry out timely backups of all sensitive and confidential data to facilitate quick recovery in case of disasters. Auditing and monitoring are important functions of a DBaaS but generate their own security risks when delegated to TPAs. Conventional methods for auditing and monitoring demand detailed knowledge of the network infrastructure and physical devices. Data privacy laws can be violated because consumers are unaware where the data is actually stored. Privacy laws in Europe and South America prohibit storing data outside the country of origin.

In summary, DBaaS data *availability* is affected by several threats including: (i) resource exhaustion caused by imprecise specification of user needs or incorrect evaluation of user specifications; and (ii) failures of the consistency management; multiple hardware and/or software failures lead to inconsistent views of user data; and failure of the monitoring and auditing system. DBaaS data *confidentiality* is affected by insider and outsider attacks, access control issues, illegal data recovery from storage, network breaches, third-party access, and inability to establish the provenance of the data.

8.8 Operating system security

An operating system allows multiple applications to share the hardware resources of a physical system subject to a set of policies. A critical function of an OS is to protect applications against a wide range of malicious attacks such as unauthorized access to privileged information, tampering with executable code, and spoofing. Such attacks can now target even single-user systems such as personal computers, tablets, or smartphones. Data brought into the system may contain malicious code; this could be the case of a Java applet, or of data imported by a browser from a malicious web site.

The *mandatory security* of an OS is considered to be [296]: “any security policy where the definition of the policy logic and the assignment of security attributes is tightly controlled by a system security policy administrator.” Access control, authentication usage, and cryptographic usage policies are all elements of the mandatory OS security.

An access control policy specifies how the OS controls the access to various system objects; authentication usage defines the authentication mechanisms used by the OS to authenticate a principal; and cryptographic usage policies specify the cryptographic mechanisms used to protect the data. A necessary but not sufficient condition for security is that the subsystems tasked to perform security-related functions are tamper-proof and cannot be bypassed. The OS should confine an application to a unique security domain.

Applications with special privileges that perform security-related functions are called *trusted applications*. Such applications should only be allowed the lowest level of privileges required to perform their functions. For example, type enforcement is a mandatory security mechanism that can be used to restrict a trusted application to the lowest level of privileges.

Enforcing mandatory security through mechanisms left to the user’s discretion can lead to a breach of security, sometimes due to malicious intent or in other cases due to carelessness or lack of understanding. Discretionary mechanisms place the burden of security on individual users. Moreover, an application may change a carefully defined discretionary policy without the consent of the user, while a mandatory policy can only be changed by a system administrator.

Unfortunately, commercial operating systems do not support multilayered security. They only distinguish between a completely privileged security domain and a completely unprivileged one. Some operating systems, e.g., Windows NT, allow a program to inherit all the privileges of the program invoking it, regardless of the level of trust in that program.

The existence of *trusted paths*, mechanisms supporting user interactions with trusted software, is critical for system security. Malicious software can impersonate trusted software when such mechanisms do not exist. Some systems allow servers to authenticate their clients and provide trusted paths for login authentication and password changing.

The solution discussed in [296] is to decompose a complex mechanism in several components with well-defined roles. For example, the access control mechanism for the application space could consist of *enforcer* and *decider* components. To access a protected object, the enforcer will gather the required information about the agent attempting the access and will pass this information to the decider together with the information about the object and the elements of the policy decision; finally, it will carry out the actions requested by the decider.

A trusted-path mechanism is required to prevent malicious software invoked by an authorized application to tamper with the attributes of the object and/or with the policy rules. A trusted path is also required to prevent an impostor from impersonating the decider agent. A similar solution is proposed

for the cryptography usage that should be decomposed into an analysis of the invocation mechanisms and an analysis of the cryptographic mechanism.

Another question is how an OS can protect itself and the applications running under it from malicious mobile code attempting to gain access to the data and the other resources and compromise system confidentiality and/or integrity. Java Security Manager uses the type-safety attributes of Java to prevent unauthorized actions of an application running in a “sandbox.” Yet, the Java Virtual Machine (JVM) accepts byte code in violation of language semantics; moreover, it cannot protect itself from tampering from other applications.

Even if all these security problems could be eliminated, the security relies on the ability of the file system to preserve the integrity of Java class code. The approach to requiring digitally signed applets and accepting them only from trusted sources could fail due to the all-or-nothing security model. A solution to securing mobile communications could be confining a browser to a distinct security domain.

Specialized *closed-box platforms*, such as the ones on some cellular phones, game consoles and Automatic Teller Machines (ATMs), could have embedded cryptographic keys that allow themselves to reveal their true identity to remote systems and authenticate the software running on them. Such facilities are not available to *open-box platforms*, the traditional hardware designed for commodity operating systems.

A highly secure operating system is necessary but not sufficient. Application-specific security is also necessary. Sometimes, security implemented above the operating system is better, e.g., electronic commerce requires a digital signature on each transaction.

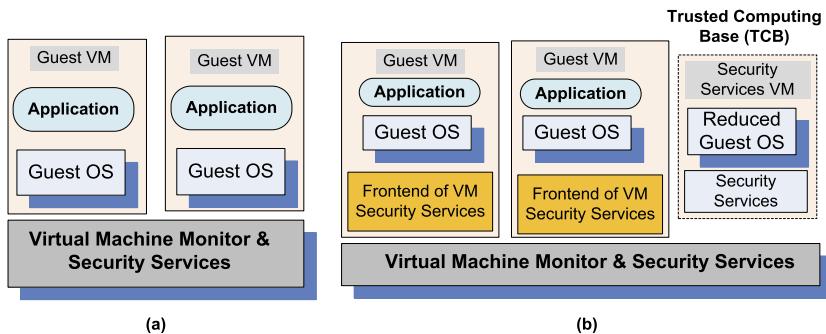
We conclude that commodity operating systems offer low assurance. Indeed, an OS is a complex software system consisting of millions of lines of code, and it is vulnerable to a wide range of malicious attacks. An OS poorly isolates one application from another; once an application has been compromised, the entire physical platform and all applications running on it can be affected. Thus, the platform security level is reduced to the security level of the most vulnerable application running on the platform.

Operating systems provide only weak mechanisms for applications to authenticate one another and do not have a trusted path between users and applications. These shortcomings add to the challenges of providing security in a distributed computing environment. For example, a financial application cannot determine if a request comes from an authorized user or from a malicious program; in turn, a human user cannot distinguish a response from a malicious program impersonating the service from the response provided by the service.

8.9 Virtual machine security

Our discussion of VM security is restricted to the traditional system VM model in Fig. 5.1(b) when the hypervisor controls the access to the hardware. The hybrid and the hosted VM models shown in Figs. 5.1(c) and (d), respectively, expose the entire system to the vulnerability of the host operating system, thus we will not analyze these models.

Virtual security services are typically provided by the hypervisor as shown in Fig. 8.3(a); another alternative is to have a dedicated VM providing security service as in Fig. 8.3(b). A secure TCB (Trusted Computing Base) is a necessary condition for security in a VM environment. When TCB is compromised, the security of the entire system is affected.

**FIGURE 8.3**

(a) Virtual security services provided by the hypervisor/Virtual Machine Monitor; (b) A dedicated security VM.

The analysis of Xen and *vBlades* in Sections 5.8 and 5.12 shows that the VM technology provides a stricter isolation of VMs from one another than the isolation of processes in a traditional operating system. Indeed, a hypervisor controls the execution of privileged operations and can thus enforce memory isolation, as well as disk and network access.

Hypervisors are less complex and better structured than traditional operating systems and in a better position to respond to security attacks. A major challenge is that a hypervisor sees only raw data regarding the state of a guest OS, while security services typically operate at a higher logical level, e.g., at the level of a file rather than a disk block.

A guest OS runs on simulated hardware, and the hypervisor has access to the state of all VMs operating on the same hardware. The state of a guest VM can be saved, restored, cloned, and encrypted by the hypervisor. Replication can ensure not only reliability but also support security, while cloning could be used to recognize a malicious application by testing it on a cloned system and observing if it behaves normally.

We can also clone a running system and examine the effect of potentially dangerous applications. Another interesting possibility is to have the guest VM's files moved to a dedicated VM and, thus, protect it from attacks [541]. This solution is possible because inter-VM communication is faster than communication between two physical machines.

Sophisticated attackers are able to fingerprint VMs and avoid VM honey pots designed to study the methods of attack. They can also attempt to access VM-logging files and, thus, recover sensitive data; such files have to be very carefully protected to prevent unauthorized access to cryptographic keys and other sensitive data.

We expect to pay some price for the better security provided by virtualization. This price includes: higher hardware costs because a virtual system requires more resources such as CPU cycles, memory, disk, and network bandwidth; the cost of developing hypervisors and modifying the host operating systems in the case of paravirtualization; and the overhead of virtualization since the hypervisor is involved in privileged operations.

VM-based intrusion detection systems such as Livewire and Siren that exploit the three capabilities of a VM for intrusion detection, isolation, inspections, and interposition are surveyed in [541]. Resource isolation was examined in Section 4.11. Inspection means that the hypervisor has the ability to review

the state of the guest VMs, and interposition means that hypervisor can trap and emulate the privileged instruction issued by the guest VMs. [541] also discusses VM-based intrusion prevention systems, such as SVFS, NetTop, and IntroVirt.

NIST security group distinguishes two groups of threats, hypervisor-based and VM-based. There are several types of hypervisor-based threats:

- a.** Resources starvation and denial of service for some VMs. Probable causes: (1) badly configured resource limits for some VMs; (2) a rogue VM with the capability of bypassing resource limits set in hypervisor.
- b.** VM side-channel attacks: malicious attack on one or more VMs by a rogue VM under the same hypervisor. Probable causes: (1) lack of proper isolation of inter-VM traffic due to misconfiguration of the virtual network residing in the hypervisor; (2) limitation of packet-inspection devices to handle high-speed traffic, e.g., video traffic; (3) presence of VM instances built from insecure VM images, e.g., a VM image having a guest OS without the latest patches.
- c.** Buffer overflow attacks.

There are two types of VM-based threats: (1) deployment of rogue or insecure VM; unauthorized users may create insecure instances from images or may perform unauthorized administrative actions on existing VMs. Probable cause: improper configuration of access controls on VM administrative tasks such as instance creation, launching, suspension, re-activation, etc.; (2) Presence of insecure and tampered VM images in the VM image repository. Probable causes: (i) lack of access control to the VM image repository; (ii) lack of mechanisms to verify the integrity of the images, e.g., digitally signed image.

8.10 Security of virtualization

The complex relationship between virtualization and security has two distinct aspects: virtualization of security and security of virtualization [303]. In Chapter 5, we praised the virtues of virtualization. We also discussed two problems associated with virtual environments: (a) the negative effect on performance due to the additional overhead; and (b) the need for more powerful systems to run multiple VMs. In this section, we take a closer look at the security of virtualization.

The complete state of an operating system running under a VM is captured by the VM. The VM state can be saved in a file, and then the file can be copied and shared. There are several useful implications of this important virtue of virtualization:

- 1.** Supports the IaaS delivery model. An IaaS user selects an image matching the local application environment and then uploads and runs the application on the cloud using this image.
- 2.** Increased reliability. An operating system with all the applications running under it can be replicated and switched to a hot standby in case of a system failure. Recall that a hot standby is a method to achieve redundancy. The primary and the backup systems run simultaneously and have identical state information.
- 3.** Straightforward mechanisms for implementing resource management policies. An OS and the applications running under it can be moved to another server to balance the load of a system. For

- example, the load of lightly loaded servers can be moved to other servers and then lightly loaded servers can be switched off or placed in standby mode to reduce power consumption.
- 4. Improved intrusion detection. In a virtual environment a clone can look for known patterns in system activity and detect intrusion. The operator can switch a server to hot standby when suspicious events are detected.
 - 5. Secure logging and intrusion protection. When implemented at the OS level, intrusion detection can be disabled and logging can be modified by an intruder. When implemented at the hypervisor layer, the services cannot be disabled or modified. In addition, the hypervisor may be able to log only events of interest for a post-attack analysis.
 - 6. More efficient and flexible software maintenance and testing. Virtualization allows the multitude of OS instances to share a small number of physical systems, instead of a large number of dedicated systems running under different operating systems, different versions of each OS, and different patches for each version.

Is there a price to pay for the benefits of virtualization? There is always the other side of a coin, so we should not be surprised that the answer to this question is a resounding Yes. In a 2005 paper [189], Garfinkel and Rosenblum argue that the serious implications of virtualization on system security cannot be ignored. This theme is revisited in 2008 by Price [404], who reached similar conclusions.

A first type of undesirable effects of virtualization lead to a diminished ability of an organization to manage its systems and track their status:

- i. The number of physical systems in the inventory of an organization is limited by cost, space, energy consumption, and human support. The explosion of the number of VMs is a fact of life; to create a VM, one simply copies a file. The only limitation for the number of VMs is the amount of storage space available.
- ii. In addition to a quantity, there is also a qualitative side to the explosion of the number of VMs. Traditionally, organizations install and maintain the same version of system software. In a virtual environment such a uniformity cannot be enforced, and the number of different operating systems, their versions, and the patch status of each version will be very diverse. This diversity will tax the support team.
- iii. One of the most critical problems posed by virtualization is related to the software lifecycle. The traditional assumption is that the software lifecycle is a straight line, hence the patch management is based on a monotonic forward progress. The virtual execution model *maps to a tree structure* rather than a line. Indeed, at any point in time, multiple instances of the VM can be created and then each one of them can be updated, different patches installed, and so on. This problem has serious implications for security, as we shall soon see.

What are the direct implications of virtualization on security? A first question is: How can the support team deal with the consequences of an attack in a virtual environment? Also, do we expect the infection with a computer virus or a worm to be less manageable in a virtual environment? The surprising answer is that an infection may last indefinitely.

Infected VMs may be dormant at the time when the measures to clean up the systems are taken. Then, at a later time, the infected VMs wake up and infect other systems. The scenario can repeat itself and guarantee that infection will last indefinitely. This is in stark contrast with the manner an infection is treated in non-virtual environments. Once an infection is detected, the infected systems are quarantined and then cleaned up; the systems will then behave normally until the next infection occurs.

A more general observation is that in a traditional computing environment a steady state can be reached. In this steady state all systems are brought up to a “desirable” state, whereas “undesirable” states, states when some of the systems are either infected by a virus or display an undesirable pattern of behavior, are only transient. The desirable state is reached by installing the latest system software version and then applying the latest patches to all systems.

A virtual environment may never reach a steady state due to the lack of control. In a nonvirtual environment, the security can be compromised when an infected laptop is connected to the network protected by a firewall, or when a virus is brought in on a removable media. But, unlike a virtual environment, the system can still reach a steady state.

A side effect of the ability to record the complete state of a VM in a file is the possibility to roll back a VM. This opens wide the door for a new type of vulnerability caused by events recorded in the memory of an attacker. Two such situations are discussed in [189].

The first is that one-time passwords are transmitted in the clear, and the protection is guaranteed only if the attacker does not have the possibility to access passwords used in previous sessions. An attacker can replay rolled-back versions and access past sniffed passwords if a system runs the S/KEY password system. S/KEY is a password system based on Leslie Lamport’s scheme. It is used by several operating systems including, Linux, OpenBSD, and NetBSD. The real password of the user is combined with a short set of characters and a counter that is decremented at each use to form a single-use password.

The second is related to the requirement of some cryptographic protocols and even non-cryptographic protocols regarding the “freshness” of the random number source used for session keys and nonces. This occurs when a VM is rolled back to a state when a random number has been generated but not yet used. A nonce is a random or pseudorandom number issued in an authentication protocol to ensure that old communications cannot be reused in replay attacks. For example, nonces are used to calculate an *MD5* of the password for *HTTP* digest access authentication. Each time the authentication challenge response code is presented, the nonces are different, thus replay attacks are virtually impossible. This guarantees that an online order to Amazon or other online store cannot be replayed.

Even noncryptographic use of random numbers may be affected by the roll-back scenario. For example, the initial sequence number for a new TCP connection must be “fresh.” The door to TCP hijacking is left open when the initial sequence number is not fresh.

Another undesirable effect of virtual environment affects trust. Recall from Section 8.5 that *trust is conditioned by the ability to guarantee the identity of entities involved*. Each computer system in a network has a unique physical, or MAC address. The uniqueness of the MAC address guarantees that an infected or a malicious system can be identified and then shut down, or denied network access. This process breaks down for virtual systems when VMs are created dynamically. Often, a random MAC address is assigned to a newly created VM to avoid name collision. The other effect discussed at length in Section 8.11 is that popular VM images are shared by many users.

The ability to guarantee confidentiality of sensitive data is yet another pillar of security affected by virtualization. Virtualization undermines the basic principle that time-sensitive data stored on any system should be reduced to a minimum. First, the owner has very limited control over where sensitive data is stored: It could be spread across many servers and may be left on some of them indefinitely. A hypervisor records the state of a VM to be able to roll it back; this process allows an attacker to access sensitive data the owner attempted to destroy.

8.11 Security risks posed by shared images

Even if we assume that a cloud service provider is trustworthy, there are other sources of concern many users either ignore, or they underestimate the danger they pose. One of them, especially critical for the IaaS cloud delivery model, is image sharing. For example, a user of AWS has the option to choose between Amazon Machine Images (AMIs) accessible through the Quick Start or the Community AMI menus of the EC2 service. The option of using one of these AMIs is especially tempting for a first-time user, or for a less sophisticated one.

First, we review the process to create an AMI. We can start from a running system, from another AMI, or from the image of a VM and copy the contents of the file system to the S3, so-called *bundling*. The first of the three steps of bundling is to create an image, the second step is to compress and encrypt the image, and the last step is to split the image into several segments and then upload the segments to the S3 storage.

Two procedures *ec2-bundle-image* and *ec2-bundle-volume* are used for the creation of an AMI. The first is used for images prepared as loopback files⁴ when the data is transferred to the image in blocks. To bundle a running system, the creator of the image can use the second procedure when bundling works at the level of the file system and files are copied recursively to the image.

To use an image, a user has to specify the resources, provide the credentials for login, a firewall configuration, and specify the region, as discussed in Section 2.2. Once the image is instantiated, the user is informed about the public DNS, and the VM is available. A Linux system can be accessed using ssh at port 22, while the Remote Desktop at port 3389 is used for Windows.

The results of an analysis carried over a period of several months, from November 2010 to May 2011, of over 5 000 AMIs available through the public catalog at Amazon are reported in [44]. Many images analyzed allowed a user to *undelete* files, recover credentials, private keys, or other types of sensitive information with little effort and using standard tools. The results of this study were shared with the Amazon's Security Team, which acted promptly to reduce the threats posed to AWS users.

The details of the testing methodology can be found in [44], but here, we only discuss the results. The study was able to audit some 5 303 images out of the 8 448 Linux AMIs and 1 202 Windows AMIs at Amazon sites in the US, Europe, and Asia. The audit covered software vulnerabilities and security and privacy risks.

The average duration of an audit was 77 minutes for a Windows image and 21 minutes for a Linux image; the average disk space used was about 1 GB and 2.7 GB, respectively. The entire file system of a Windows AMI was audited because most of the malware targets Windows systems. Only directories containing executables for Linux AMIs were scanned; this strategy and the considerably longer start-up time of Windows explain the time discrepancy of the audits for the types of AMIs.

A *software vulnerability* audit revealed that 98% of the Windows AMIs (249 out of 253) and 58% (2 005 out of 3 432) Linux AMIs audited had critical vulnerabilities. The average number of vulnerabilities per AMI were 46 for Windows and 11 for Linux AMIs. Some of the images were rather old; 145,

⁴ A *loopback file system* (LOFS) is a virtual file system providing an alternate path to an existing file system. When other file systems are mounted onto an LOFS file system, the original file system does not change. One useful purpose of LOFS is to take a CDROM image file, a file of type “.iso”, mount it on the file system and then access it without the need to record a CD-R. It is somewhat equivalent to the Linux *mount -o loop* option but adds a level of abstraction; most commands that apply to a device can be used to handle the mapped file.

38, and 2 Windows AMIs and 1197, 364, and 106 Linux AMIs were older than two, three, and four years, respectively. The tool used to detect vulnerabilities, the Nessus system, available from <http://www.tenable.com/productus/nessus>, classifies the vulnerabilities based on the severity in four groups, at levels zero to three. The audit reported only vulnerabilities of the highest severity level, e.g., remote code execution.

Three types of *security risks* are analyzed: (1) backdoors and leftover credentials, (2) unsolicited connections, and (3) malware. An astounding finding, about 22% of scanned Linux AMIs contained credentials allowing an intruder to remotely login to the system. Some 100 passwords, 995 ssh keys, and 90 cases when both could be retrieved were identified.

To rent a Linux AMI, a user must provide the public part of her ssh key, and this key is stored in the *authorized_keys* in the home directory. This opens a backdoor for a malicious creator of an AMI who does not remove her own public key from the image and can remotely login to any instance of this AMI. Another backdoor is opened when the ssh server allows password-based authentication and the malicious creator of an AMI does not remove her own password. This backdoor is open even wider since one can extract the password hashes and then crack the passwords using a tool such as John the Ripper, see <http://www.openwall.com/john>.

Another threat is posed by the omission of the *cloud-init* script that should be invoked when the image is booted. This script provided by Amazon regenerates the host key an ssh server uses to identify itself; the public part of this key is used to authenticate the server. When this key is shared among several systems, these systems become vulnerable to man-in-the middle attacks.

An attacker impersonates the agents at both ends of a communication channel in the *man-in-the-middle* attack and makes them believe that they are communicating through a secure channel. For example, if B sends her public key to A, but C is able to intercept it, such an attack proceeds as follows: C sends a forged message to A claiming to be from B, but instead includes C's public key. Then A encrypts her message with C's key, believing that she is using B's key, and sends the encrypted message to B. The intruder, C, intercepts, deciphers the message using her private key, possibly alters the message, and re-encrypts with the public key B originally sent to A. When B receives the newly encrypted message, she believes it came from A.

When this script does not run, an attacker can use the *NMap* tool⁵ to match the *ssh* keys discovered in the AMI images with the keys obtained with *NMap*. The study reports that the authors were able to identify more than 2 100 instances following this procedure.

Unsolicited connections pose a serious threat to a system. Outgoing connections allow an outside entity to receive privileged information, e.g., the IP address of an instance and events recorded by a *syslog* daemon to files in the *var/log* directory of a Linux system. Such information is available only to users with administrative privileges.

The audit detected two Linux instances with modified *syslog* daemon that forwarded to an outside agent information about events such as login and incoming requests to a web server. Some of the unsolicited connections are legitimate, for example, connections to a software update site. It is next to impossible to distinguish legitimate from malicious connections.

⁵ *NMap* is a security tool running on most operating systems including *Linux*, *Microsoft Windows*, *Solaris*, *HP-UX*, *SGI-IRIX*, and *BSD* variants, such as *Mac OS X* to map the network. Mapping the network means to discover hosts and services in a network.

Malware including viruses, worms, spyware, and Trojans were identified using *ClamAV*, a software tool with a database of some 850 000 malware signatures, available from <http://www.clamav.net>. Two infected Windows AMIs were discovered, one with a *Trojan-Spy* (variant 50112) and a second one with a *Trojan-Agent* (variant 173287). The first Trojan carries out key logging and allows stealing data from the files system and monitoring processes; the AMI also included a tool to decrypt and recover passwords stored by the *Firefox* browser, called *Trojan.Firepass*.

The creator of a shared AMI assumes some *privacy risks*. Her private keys, IP addresses, browser history, shell history, and deleted files can be recovered from the published images. A malicious agent can recover the AWS API keys that are not password protected. Then the malicious agent can start AMIs and run cloud applications at no cost to herself because the computing charges are passed on to the owner of the API key. The search can target files with names such as *pk-[0-9A-Z]*.pem* or *cert-[0-9A-Z]*.pem* used to store API keys.

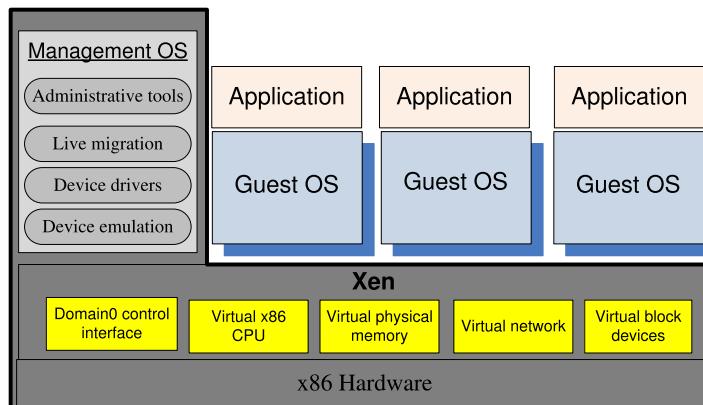
Another avenue for a malicious agent is to recover ssh keys stored in files named *id_dsa* and *id_rsa*. Though ssh keys can be protected by a passphrase, the audit determined that the majority of ssh keys (54 out of 56) were not password protected. A *passphrase* is a sequence of words used to control access to a computer system. A passphrase is the analog of a password but provides added security. For high-security nonmilitary applications, NIST recommends an 80-bit strength passphrase. Therefore, a secure passphrase should consist of at least 58 characters including uppercase and alphanumeric characters. The entropy of written English is less than 1.1 bits per character.

Recovery of IP addresses of other systems owned by the same user requires access to the *lastlog* or the *lastb* databases. The audit found 187 AMIs with a total of more than 66 000 entries in their *lastb* databases. Nine AMIs contained Firefox browser history and allowed the auditor identify the domains contacted by the user.

Six hundred and twelve AMIs contained at least one shell history file. The audit analyzed 869 history files named *~/.history*, *~/.bash_history*, and *~/.sh_history* containing some 160 000 lines of command history and identified 74 identification credentials. The users should be aware that, when the HTTP protocol is used to transfer information from a user to a web site, the GET requests are stored in the logs of the web server. Passwords and credit card numbers communicated via a GET request can be exploited by a malicious agent with access to such logs. When remote credentials, such as the DNS management password, are available, then a malicious agent can redirect traffic from its original destination to her own system.

Recovery of deleted files containing sensitive information poses another risk. When the sectors on the disk containing sensitive information are actually overwritten by another file, recovery of sensitive information is much harder. To be safe, the creator of the image effort should use utilities, such as *shred*, *scrub*, *zerofree* or *wipe*, to make recovery of sensitive information next to impossible. If the image is created with the block-level tool discussed at the beginning of this section, the image will contain blocks of the file system marked as free; such block may contain information from deleted files. The audit process was able to recover files from 98% of the AMIs using the *exundelete* utility. The number of files recovered from an AMI were as low as 6 and as high as 40 000.

We conclude that the users of published AMIs, as well as the providers of images, may be vulnerable to a wide range of security risks and must be fully aware of the dangers posed by image sharing.

**FIGURE 8.4**

The trusted computing base of a Xen-based environment includes the hardware, Xen, and the management operating system running in *Dom0*. The management OS supports administrative tools, live migration, device drivers, and device emulators. A guest OS and applications running under it reside in a *DomU*.

8.12 Security risks posed by a management OS

We often hear that virtualization enhances security because a VM hypervisor is considerably smaller than an operating system. For example, the Xen hypervisor discussed in Section 5.8 has approximately 60 000 lines of code, one to two orders of magnitude fewer than a traditional operating system.⁶

A hypervisor supports a stronger isolation between the VMs running under it than the isolation between processes supported by a traditional operating system. Yet the hypervisor must rely on a management OS to create VMs and to transfer data in and out from a guest VM to storage devices and network interfaces.

A small hypervisor can be carefully analyzed, thus one could conclude that the security risks in a virtual environment are diminished. We have to be cautious with such sweeping statements. Indeed, the Trusted Computer Base (TCB)⁷ of a cloud computing environment includes not only the hypervisor but also the management OS. The management OS supports administrative tools, live migration, device drivers, and device emulators.

For example, the TCB of an environment based on Xen includes not only the hardware and the hypervisor, but also the management operating system running in *Dom0*; see Fig. 8.4. System vulnerabilities can be introduced by both Xen and the management operating system. An analysis of Xen vulnerabilities reports that 21 of 23 attacks were against service components of the control VM [114];

⁶ The number of lines of code of the *Linux* operating system evolved over time from 176 250 for *Linux 1.0.0*, released in March 1995 to 1 800 847 for *Linux 2.2.0*, released in January 1999, 3 377 902 for *Linux 2.4.0*, released in January 2001, and to 5 929 913 for *Linux 2.6.0*, released in December 2003.

⁷ The TCB is defined as the totality of protection mechanisms within a computer system, including hardware, firmware, and software. The combination of all these elements is responsible for enforcing a security policy.

11 attacks were attributed to problems in the guest OS caused by buffer overflow and 8 were denial of service attacks. Buffer overflow allows execution of arbitrary code in the privileged mode.

Dom0 manages building of user domains (*DomU*), a process consisting of several steps:

- a. Allocate memory in *Dom0* address space and load the kernel of guest OS from secondary storage.
- b. Allocate memory and use foreign mapping to load the kernel to the new VM. The foreign mapping mechanism of Xen is used by *Dom0* to map arbitrary memory frames of a VM into its page tables.
- c. Set up the initial page tables for the new VM.
- d. Release the foreign mapping on the new VM memory, set up the virtual CPU registers, and launch the new VM.

A malicious *Dom0* can play several nasty tricks at the time that it creates a *DomU*:

- i. Refuse to carry out the steps necessary to start the new VM, an action that can be considered a *denial-of-service* attack.
- ii. Modify the kernel of the guest OS in ways that will allow a third party to monitor and control the execution of applications running under the new VM.
- iii. Undermine the integrity of the new VM by setting the wrong page tables and/or setup wrong virtual CPU registers.
- iv. Refuse to release the foreign mapping, and access the memory while the new VM is running.

We now discuss the run-time interaction between *Dom0* and a *DomU*. Recall that *Dom0* exposes a set of abstract devices to the guest operating systems using *split drivers*; the front end of such a driver is in *DomU*, its back end in *Dom0*, and the two communicate via a ring in shared memory; see Section 5.8.

In the original implementation of Xen, a service running in a *DomU* sends data to, or receives data from, a client located outside the cloud using a network interface in *Dom0*; it transfers the data to I/O devices using a device driver in *Dom0*. Later implementations of Xen offer the pass-through option.

We have to ensure that run-time communication through *Dom0* is encrypted. Yet, Transport Layer Security (TLS) does not guarantee that *Dom0* cannot extract cryptographic keys from the memory of the OS and applications running in *DomU*. A significant security weakness of *Dom0* is that the entire state of the system is maintained by XenStore; see Section 5.8. A malicious VM can deny access to this critical element of the system to other VMs; it can also gain access to the memory of a *DomU*. This brings us to additional requirements for confidentiality and integrity imposed on *Dom0*.

Dom0 should be prohibited to use foreign mapping for sharing memory with a *DomU*, unless *DomU* initiates the procedure in response to a hypercall from *Dom0*. When this happens, *Dom0* should be provided with an encrypted copy of the memory pages and of the virtual CPU registers. The entire process should be closely monitored by the hypervisor which, after the access, should check the integrity of the affected *DomU*.

A virtualization architecture that guarantees confidentiality, integrity, and availability for the TCB of a Xen-based system is presented in [303]. A secure environment, when *Dom0* cannot be trusted, can only be ensured if the guest application is able to store, communicate and process data safely. The guest software should have access to a secure secondary storage on a remote storage server and to the network interfaces when communicating with the user. A secure run-time system is also needed.

To implement a secure run-time system, we have to intercept and control the hypercalls used for communication between a *Dom0* that cannot be trusted and a *DomU* we want to protect. Hypercalls issued by *Dom0* that do not read from or write to the memory of a *DomU* or to its virtual registers should

be allowed. Other hypercalls should be restricted either completely or during a specific time window. For example, hypercalls used by *Dom0* for debugging or for the control of the IOMMU should be prohibited. The Input/Output Memory Management Unit (IOMMU) connects the main memory with a DMA-capable I/O bus; it maps device-visible virtual addresses to physical memory addresses and provides memory protection from misbehaving devices.

We cannot restrict some of the hypercalls issued by *Dom0*, even though they can be harmful to the security of *DomU*. For example, foreign mapping and access to the virtual registers are needed to save and restore the state of *DomU*. We should check the integrity of *DomU* after the execution of such security-critical hypercalls.

New hypercalls are necessary to protect:

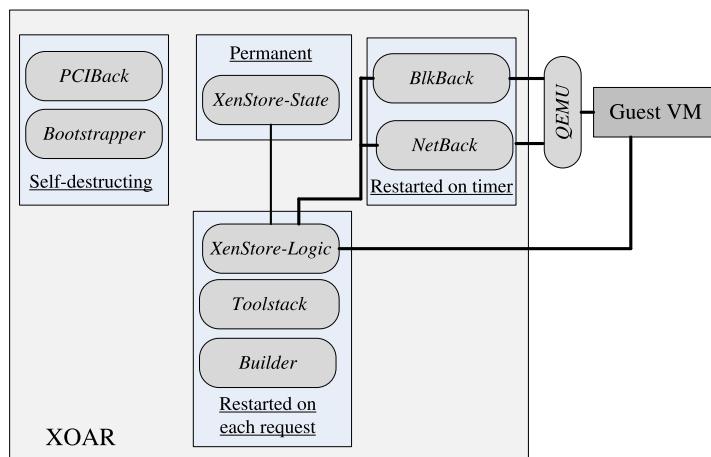
1. The privacy and integrity of the virtual CPU of a VM. When *Dom0* wants to save the state of the VM, the hypercall should be intercepted and the contents of the virtual CPU registers should be encrypted. The virtual CPU context should be decrypted, and then an integrity check should be carried out when *DomU* is restored.
2. The privacy and integrity of the VM virtual memory. The *page table update* hypercall should be intercepted and the page should be encrypted so that *Dom0* handles only encrypted pages of the VM. The hypervisor should calculate a hash of all the memory pages before they are saved by *Dom0* to guarantee the integrity of the system. Address translation is necessary because a restored *DomU* may be allocated a different memory region [303].
3. The freshness of the virtual CPU and the memory of the VM. The solution is to add to the hash a version number.

As expected, the increased level of security and privacy leads to an increased overhead. Measurements reported in [303] show increases by a factor of: 1.7 to 2.3 for the domain build time, 1.3 to 1.5 for the domain save time, and 1.7 to 1.9 for the domain restore time.

8.13 Xoar—breaking the monolithic design of the TCB

Xoar is a modified version of Xen designed to boost system security [114]. The security model of Xoar assumes that the system is professionally managed and that a privileged access to the system is granted only to system administrators. The model also assumes that the administrators have neither financial incentives, nor the desire to violate the trust of the user. The security threats come from a guest VM that could attempt to violate the data integrity or the confidentiality of another guest VM on the same platform, or to exploit the code of the guest. Other sources of threats are bugs in initialization code of the management VM. Xoar is based on micro-kernel⁸ design principles. Xoar modularity makes exposure to risk explicit and enables the guests to configure the access to services based on their needs. Modularity allows the designers of Xoar to reduce the size of the permanent footprint of the system and

⁸ A microkernel (μ -kernel) supports only the basic functionality of an operating system kernel including low-level address space management, thread management, and interprocess communication. Traditional operating system components, such as device drivers, protocol stacks, and file systems, are removed from the microkernel and run in the user space.

**FIGURE 8.5**

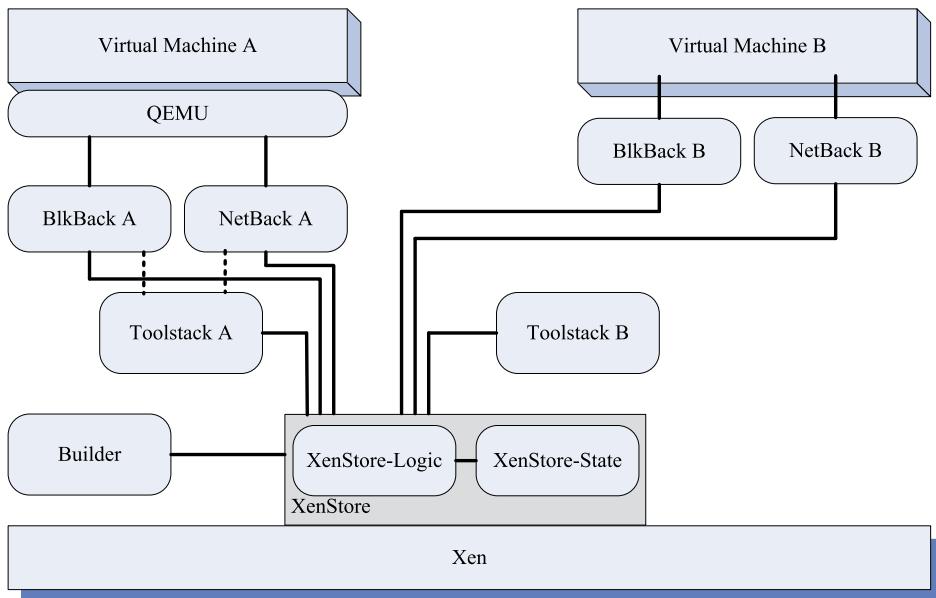
Xoar has nine classes of components of four types: permanent, self-destructing, restarted upon request, and restarted on timer. A guest VM is started using the *Toolstack* by the *Builder* and is controlled by the *XenStore-Logic*. The devices used by the guest VM are emulated by the *QEMU* component.

increase the level of security of critical components. The ability to record a secure audit log is another critical function of a hypervisor facilitated by a modular design. The design goals of Xoar are:

- Maintain the functionality provided by Xen.
- Ensure transparency with existing management and VM interfaces.
- Tight control of privileges. Each component should only have the privileges required by its function.
- Minimize the interfaces of all components to reduce the possibility that a component can be used by an attacker.
- Eliminate sharing. Make sharing explicit whenever it cannot be eliminated to allow meaningful logging and auditing.
- Reduce the opportunity of an attack targeting a system component by limiting the time window during which the component runs.

These design principles aim to break the monolithic TCB design of a *Xen*-based system. Inevitably, this strategy has an impact on performance, but the implementation should attempt to keep the modularization overhead to a minimum.

A close analysis shows that booting the system is a complex activity, but the fairly large modules used during booting are no longer needed once the system is up and running. In Section 5.8, we have seen that *XenStore* is a critical system component because it maintains the state of the system, and, thus, it is a prime candidate for hardening. The *ToolStack* is only used for management functions and can only be loaded upon request. Xoar has four types of components: permanent, self-destructing, restarted upon request, and restarted on timer; see Fig. 8.5:

**FIGURE 8.6**

Component sharing between guest VM in Xoar. Two VMs share only the *XenStore* components. Each one has a private version of the *BlkBack*, *NetBack*, and *Toolstack*.

1. Permanent components. *XenStore-State* maintains all information regarding the state of the system.
2. Components used to boot the system; they self-destruct before any user VM is started. The two components discover the hardware configuration of the server including the PCI drivers and then boot the system: *PCIBack*—virtualizes access to PCI bus configuration and *Bootstrapper*—coordinates booting of the system.
3. Components restarted on each request: *XenStore-Logic*; *Toolstack*—handles VM management requests, e.g., it requests the *Builder* to create a new guest VM in response to a user request; and *Builder*—initiates user VMs.
4. Components restarted on a timer: the two components export physical storage device drivers and the physical network driver to a guest VM. *Blk-Back*—exports physical storage device drivers using *udev*⁹ rules. *NetBack*—exports the physical network driver.

Another component, QEMU, is responsible for device emulation. *Bootstrapper*, *PCIBack*, and *Builder* are the most privileged components, but the first two are destroyed once Xoar is initialized. The *Builder* is very small, it consists of only 13 000 lines of code. XenStore is broken into two components, *XenStore-Logic* and *XenStore-State*. Access control checks are done by a small monitor module in *XenStore-State*. Guest VMs share only the *Builder*, *XenStore-Logic*, and *XenStore-State*; see Fig. 8.6.

⁹ *udev* is the device manager for the Linux kernel.

Users of Xoar are able to share only service VMs with guest VMs that they control; to do so, they specify a tag on all of the devices of their hosted VMs. Auditing is more secure: whenever a VM is created, deleted, stopped, or restarted by Xoar, the action is recorded in an append-only database on a different server accessible via a secure channel.

Rebooting provides the means to ensure that a VM is in a known good state. To reduce the overhead and the increased startup time demanded by a reboot, Xoar uses *snapshots* instead of rebooting. The service VM snapshots itself when it is ready to service a request; similarly, snapshots of all components are taken immediately after their initialization and before they start interacting with other services or guest VMs. Snapshots are implemented using a copy-on-write (COW) mechanism¹⁰ to preserve any page about to be modified.

8.14 Mobile devices and cloud security

Mobile devices are an integral part of the cloud ecosystem; mobile applications use cloud services to access and store data or to carry out a multitude of computational tasks. The security challenges for mobile devices are common to all computer and communication systems and include: (i) Confidentiality—ensure that transmitted and stored data cannot be read by unauthorized parties; (ii) Integrity—detect intentional or unintentional changes to transmitted and stored data; (iii) Availability—ensure that users can access cloud resources whenever needed; and (iv) Nonrepudiation—the ability to ensure that a party to a contract cannot deny the sending of a message that they originated.

The technology stack of a mobile device consists of the hardware, the firmware, the operating system, and the applications. The separation between the firmware and the hardware of a mobile device is blurred. A baseband processor is used solely for telephony services involving data transfers over cellular networks operating outside the control of the mobile OS that runs on the application processor. Security-specific hardware and firmware stores encryption keys, certificates, credentials, and other sensitive information on some mobile devices.

The nature of mobile devices places them at higher exposure to threats than stationary ones. Mobile devices are designed to easily install applications, to use third-party applications from application stores, and to communicate with computer clouds via often untrusted cellular and WiFi networks. Mobile devices interact frequently with other systems to exchange data and often use untrusted content.

Mobile devices often require a short authentication passcode and may not support strong storage encryption. Location services increase the risk of targeted attacks. Potential attackers are able to determine user's location, correlate the location with information from other sources on the individuals the user associates with, and infer other sensitive information. Special precautions must then be taken due to exposure to the unique security threats affecting mobile devices, including: (i) mobile malware; (ii) stolen data due to loss, theft, or disposal; (iii) unauthorized access; (iv) electronic eavesdropping and electronic tracking; and (v) access to data by third-party applications. Some of these threats can

¹⁰ Copy-on-write (COW) is used by virtual memory operating systems to minimize the overhead of copying the virtual memory of a process when a process creates a copy of itself. Then, the pages in memory that might be modified by the process or by its copy are marked as COW. When one process modifies the memory, the operating system's kernel intercepts the operation and copies the memory so that changes in one process's memory are not visible to the other.

propagate to the cloud infrastructure a mobile device is connected to. For example, files stored on the mobile devices subject to ransomware and encrypted by a malicious intruder can migrate to the backup stored on the cloud. The risks posed to the cloud infrastructure by mobile devices are centered around data leakage and compromise. Such security risks are due to a set of reasons including:

1. Loss of the mobile device, lock-screen protection, enabling smudge attacks, and other causes leading to mobile access control. A smudge attack is a method to discern the password pattern of a touchscreen device such as a cell phone or tablet computer.
2. Lack of confidentiality protection for data in transit in unsafe or untrusted WiFi or cellular networks.
3. Unmatched firmware or software including operating system and application software bypassing the security architecture, e.g., rooted/jailbroken devices.
4. Malicious mobile applications bypassing access control mechanisms.
5. Misuse or misconfiguration of location services, such as GPS.
6. Acceptance of fake mobility management profiles.

An in-depth discussion of the Enterprise Mobile Management (EMM) lists various EMM services, including the Mobile Device Management (MDM) and the Mobile Application Management (MAM), and suggests a number of functional and security capabilities of the system [184]. Some of these policies and mechanisms should also be applied to mobile devices connected to computer cloud:

- i. Use device encryption, application-level encryption, and remote wipe capabilities to protect storage.
- ii. Use Transport Layer Security (TLS) for all communication channels.
- iii. Isolate user-level applications from each other to prevent data leakage between applications using sandboxing.
- iv. Use device integrity checks for boot validation, verified application, and OS updates.
- v. Use auditing and logging.
- vi. Enforce authentication of the device owner.
- vii. Automatic, regular device integrity and compliance checks for threats and compliance.
- viii. Automated alerts for policy violations.

A system for Microsoft Outlook mobile application requires individuals who wish to participate in a managed scenario to download the Microsoft Community Portal application and input the required information including local authentication to the mobile OS via a lock screen and the encryption capabilities provided by the mobile OS to protect data on the device. The cloud MDM portal discussed in [184] is available to administrators through a web interface.

8.15 Mitigating cloud vulnerabilities in the age of ransomware

In Section 2.13, we have seen that computer clouds are changing the enterprise computing landscape at stunning speed. The massive migration to cloud computing is motivated by the desire to reduce costs and operate in a more secure environment. The latter rationale is being challenged by this explosive migration; the number of potentially vulnerable and attractive targets increases and hackers are becoming more sophisticated and able to exploit vulnerabilities in cloud systems.

Data breaches due to cloud misconfigurations cost businesses nearly \$3.18 trillion in 2019. Some vulnerabilities are exposed by the increasing larger size of the cloud user population. Gartner predicts that by 2025, 99% of cloud security failures will be due to customers fault.

Nearly two-thirds of businesses and organizations view cloud security as the biggest impediment to adoption. According to CSA, the top security issues affecting organizations using the cloud are: malware proliferation (63%), advanced persistent threats, (53%) compromised accounts (43%), and insider threats (42%). We start the discussion of cloud vulnerabilities in 2021 with a few relevant statistics regarding enterprises on the cloud¹¹:

1. 29% of enterprises on the cloud have experienced potential account compromises.
2. 27% of enterprises on the cloud permit root user activities.
3. 49% of enterprises on the cloud do not encrypt their cloud databases.
4. 21% of cloud-hosted files include sensitive data, a 17% increase in the past two years.
5. 83% of businesses worldwide store sensitive data on the cloud.
6. 41% of access keys remain unchanged on the cloud in the past 90 days.
7. 22% of cloud users share files, and as much as 48% of cloud files eventually get shared.
8. 20% of the sensitive data passing through clouds is due to email services; this volume has risen by 59% in the past two years.

The most significant cloud security threats are: (a) access management; (b) data breaches and data leaks; (c) data loss; (d) insecure APIs; and (e) misconfigured cloud storage.

Access management enforces a rigorous access policy and uses authentication and identity verification tools. To mitigate the threats, many CSPs now offer multifactor authentication systems as part of their standard packages. At the same time, IT administrators in charge of access management need to periodically audit the level of access of all employees and promptly remove the privileges of departing employees; they should deploy the most secure authentication and identity verification tools available.

Data breaches are a significant threat in cloud systems. Attackers take advantage of expired digital certificates, as in the case of *Equifax* when personal data of more than 148 million individuals was stolen by hackers in 2017. To mitigate this threat, organizations should use encryption for email servers and messages. Organizations should also use digital certificates, such as SSL/TLS website certificates and S/MIME (secure/multipurpose Internet mail extension) certificates.

Data loss is a major concern for cloud users. An average of 51% of organizations have publicly exposed at least one cloud storage service. 84% of organizations report that traditional security solutions do not work in cloud environments. Regular and thorough backups on secure storage devices are critical to protect valuable data from ransomware attacks. A hacker with access to the systems used to access the cloud could encrypt user data files that are then automatically backed up to the cloud; a hacker may also encrypt cloud storage and demand payment for returning data.

Many APIs still have security vulnerabilities, e.g., some APIs give CSPs access to user data. It was recently revealed that Facebook and Google stored user passwords in plaintext, allowing staff to access them. Before moving to a CSP any organization should make sure that the CSP adheres to OWASP API security guidelines see <https://owasp.org/www-project-api-security/>.

¹¹ The statistics and the data on security threats are from <https://cloud-standards.org/cloud-computing-statistics/> and <https://securityboulevard.com/2020/05/cloud-security-5-serious-emerging-cloud-computing-threats-to-avoid/>, respectively.

Misconfiguration of cloud storage exposes confidential data. This can be caused by either data stored in large and confusing structures leaving important files unprotected, or by failing to change the default security settings on the cloud storage. An example of misconfigured cloud storage is the National Security Agency mistake that made a number of top-secret documents available via an external browser.

The shared security responsibility (SSR) model discussed in Section 2.1 and illustrated by Table 2.1 is a powerful concept for cloud security with different implications for different cloud delivery models. It may seem counterintuitive that SaaS is a most vulnerable target though the cloud infrastructure and the software stack are all controlled by the CSP.

A closer look reveals that SaaS access to data is solely the responsibility of the large number of users, some lacking basic elements of computer literacy. Even the most sophisticated SaaS uses have minimal control or visibility of the underlying infrastructure of the services they use. Malicious attacks such as *GoldenEye* and *XcodeGhost* ransomware show that hackers are focused on high-value targets—such as health care institutions and critical services provided by local communities.

The most relevant SaaS security issues are: (i) the lack of total control over who can access sensitive data in the cloud; (ii) the inability to control and keep track of data in transit to and from cloud applications; (iii) the inability to prevent the misuse of data or theft by insiders; (iv) the lack of skilled staff to manage the security of cloud applications used by an organization; and (v) the inability to assess the security operations of the cloud service providers.

Any attack on the IaaS infrastructure affects a large number of users. IaaS user responsibilities reflect the extended degrees of freedom allowed to IaaS users. A hostile takeover of resources allocated to a user and their use as bots to launch attacks against third parties by malicious individuals is possible. Advanced persistent attacks mounted on the IaaS infrastructure are possible.

8.16 AWS security

To conclude this chapter, we take a closer look at the security support available at one IaaS cloud service provider. We choose AWS due to the large number of organizations that demand high security standards hosted by Amazon. Security experts monitor the vast AWS infrastructure and build and maintain security services to prevent, detect, respond, and remediate the effects of security threats. We start with a review of AWS security tools.¹²

CloudTrail is a service that enables governance, compliance, operational auditing, and risk auditing of an AWS account. The event history, including Management Console activity, AWS SDK (Software Development Kit), and all AWS services invoked, simplifies security analysis, resource change tracking, and troubleshooting.

The Inspector (IS) is an automated security assessment service; it assesses applications for exposure, vulnerabilities, and deviations from best practices. After performing an assessment, IS produces a detailed list of security findings prioritized by level of severity and checks for unintended network accessibility of EC2 instances and for vulnerabilities on those EC2 instances.

¹² An in-depth of discussion of AWS security can be found at <https://aws.amazon.com/security/>.

Web Application Firewall (WAF) protects web applications or APIs against common web exploits and bots that may compromise security, consume excessive resources, or affect availability. WAF provides control over the traffic generated by applications and blocks attack patterns, such as SQL injection¹³ or cross-site scripting (XSS).¹⁴

Cognito is used for identity management. It can detect brute-force authentication and fraudulent login attempts. Cognito User Pools supports identity and access management standards, such as Oauth 2.0, SAML 2.0, and OpenID Connect.

CloudHSM helps generate encryption keys using managed hardware security modules. An AWS user can configure the Key Management Service (KMS) to use CloudHSM cluster as a custom key store, rather than the default KMS key store. It automatically load balances requests and securely duplicates keys stored in any HSM to all of the other HSMs in a cluster.

CloudFront is a content delivery network. It protects applications from DDoS attacks and transfers data securely at high speeds. It offers security capabilities, such as field level encryption and HTTPS support, seamlessly integrated with Shield, WAF, and Route 53 to protect against multiple types of attacks including network and application layer DDoS (Distributed Denial of Service) attacks. With edge compute features and Lambda@Edge, one can run code across AWS locations globally.

We believe that some of the most consequential aspects of AWS security are:

- i. Fine-grain identity, access controls, and continuous monitoring for near real-time security information.
- ii. Automation to reduce human configuration errors.
- iii. Extended use of encryption. Data flowing across the AWS global network is automatically encrypted at the physical layer before leaving a secured facilities.
- iv. The added level of competence provided by a global program of Technology and Consulting Partners specialized in security—the AWS Partner Network.

8.17 Further readings

Cloud Security Alliance (CSA) is an organization with more than 100 corporate members. It aims to address all aspects of cloud security and serve as a cloud security standards incubator. The reports, available from the web site of the organization, are periodically updated; the original report was published in 2009 [120], and subsequent reports followed [121] and [122]. An open security architecture is presented in [380].

A seminal paper [189] on the negative implications of virtualization on system security “When virtual is harder than real: security challenges in VM based computing environments” by Garfinkel and Rosenblum was published in 2005, followed by another one that reached similar conclusions, [404]. Risk and trust are analyzed in [152], [254], and [319]. Cloud security is also discussed in [344], [459], and [465]. Managing cloud information leakage is the topic of [512].

¹³ Malicious SQL statements are inserted into an entry field for execution.

¹⁴ Malicious scripts are injected into trusted websites enabling web application to send malicious code as a browser script to a different end user.

A 2010 paper [212] presents a taxonomy of attacks on computer clouds and [137] covers the management of security services lifecycle. Security issues vary depending on the cloud model as discussed in [375]. The privacy impact in cloud computing is the topic of [469]. A 2011 book [515] gives a comprehensive look at cloud security. Privacy and protection of personal data in the EU is discussed in a document available at <http://ec.europa.eu/justice/policies/privacy>.

The paper [35] analyzes the inadequacies of current risk controls for the cloud. Intercloud security is the theme of [59]. Secure collaborations are discussed in [62]. The paper [304] presents an approach for secure VM execution under untrusted management OS. The social impact of privacy in cloud computing is analyzed in [166]. An anonymous access control scheme is presented in [260].

An empirical study into the security exposure to hosts of hostile virtualized environments can be found at <http://taviso.decsystem.org/virtsec.pdf>. A model-based security testing approach for cloud computing is presented in [535]. Cold boot attacks on encryption keys are discussed in [222]. Cloud security concerns and mobile device security are covered in several NIST documents. [401] introduces an encryption system for query processing and [435] discusses a pragmatic security discipline.

8.18 Exercises and problems

- Problem 1.** Identify the main security threats for the SaaS cloud delivery model on a public cloud. Discuss the various aspects of these threats on a public cloud as compared to the threats posed to similar services provided by a traditional service-oriented architecture running on a private infrastructure.
- Problem 2.** Analyze how the six attack surfaces discussed in Section 8.2 and illustrated in Fig. 8.1 apply to the SaaS, PaaS, and IaaS cloud delivery models.
- Problem 3.** Analyze Amazon privacy policies, and design a service level agreement you would sign if you were to process confidential data using AWS.
- Problem 4.** Analyze the implications of the lack of trusted paths in commodity operating systems and give one or more examples showing the effects of this deficiency. Analyze the implications of the two-level security model of commodity operating systems.
- Problem 5.** Compare the benefits and the potential problems due to virtualization on public, private, and hybrid clouds.
- Problem 6.** Read [44] and discuss the measures taken by Amazon to address the problems posed by shared images available from AWS. Would it be useful to have a cloud service to analyze images and sign them before being listed and made available to the general public?
- Problem 7.** Analyze the risks posed by foreign mapping and the solution adopted by Xoar. What is the security risk posed by XenStore?
- Problem 8.** Read [114] and discuss the performance of the system. What obstacles to its adoption by the providers of IaaS services can you foresee?
- Problem 9.** Discuss the impact of international agreements regarding privacy laws on cloud computing.
- Problem 10.** Propagation of the malware in the Internet has similarities with the propagation of an infectious disease. Discuss the three models for the propagation of an infectious disease in a finite population, *SI*, *SIR*, and *SIS*. Justify the formulas describing the dynamics of the system for each model. *Hint:* read [73, 160] and [270].