

Index

A

Absolute deadline, 316
Acceptor, 398
Access Control List (ACL), 151, 227, 423
Access layer, 200
Access transparency, 73
Account, 261
Action, 106, 151
Active space, 421
Activities, 416
Actual virtual time, 311
Actuators, 334
Advanced Microcontroller Bus Architecture (AMBA), 65
Advanced Research Projects Agency (ARPA), 93
Advanced RISC Machine (ARM), 159
Advisory locks, 234
Aggregate layer, 200
Aggregation, 136
Aggregation function, 458
Algorithm
 allocation, 322
 edge, 126
 generic contraction, 556
 generic TLC, 554
 network resource management, 204
 Paxos, 236, 237
 rate monotonic, 298
Algorithmic communication, 354
Allocate thread, 385
Allocation algorithms, 322
Allocation manager, 230
Amazon, 15, 18, 23, 24
Amazon Machine Image (AMI), 19, 22, 137, 278, 430
Amazon Web Services (AWS), 1, 19
Amazon Web Services Licensing Agreement (AWSLA), 24
Amdahl's Law, 68
Amplitude amplifications, 511
Analytics, 298
Andrew File System (AFS), 216, 227
Ansible, 18
Anti-entropy, 250
Anticipated exceptions, 415
Apache
 capacity scheduler, 306
 Hadoop, 105, 320, 428, 431, 432
Aperiodic tasks, 316

Application

arbitrarily divisible, 329
coarse-grained data-parallel, 104
controllers, 335
data, 2, 6
data-parallel, 99
dynamic data-driven, 463
embarrassingly parallel, 68
layer, 16, 80
modularly divisible, 329
request-parallel, 99
trusted, 272
VM, 138
Application Binary Interface (ABI), 82
Application Program Interface (API), 82, 425
Application Specific Integrated Circuit (ASIC), 63
Application/logic tier, 4
Arbitrarily divisible application, 329
Arbitrarily divisible load sharing model, 412
Architecture, 42, 48
 ARM, 55, 57
 CISC, 43, 48
 computer, 42, 140
 taxonomy, 46
 control flow, 43
 data flow, 43, 92
 hourglass, 178
 Itanium, 142, 161, 173
 load-store, 43
 network, 154, 176
 ONF, 188
 scalable data center communication, 200
 SIMD, 58–60
 software, 4
 vector, 59
 von Neumann, 91
Arithmetic intensity (ArI), 60
ARM, 450
 architecture, 55, 57
 processors, 159, 160
ARM Trust Zone (ATZ), 160
Array, 98
 switch, 98
Artificial Intelligence (AI), 3, 11, 25, 501, 503
Assured rate (AR), 207
Asymmetric confusion, 364

- Asymmetric core processor, 70
 Atlimit, 207
 Atomic transaction, 111
 Atomicity, 394
 Atomicity log, 236
 Audit, 15
 Auditor, 14
 Auto Scaling, 22
 Automatic speech recognition (ASR), 507
 Autonomic performance managers, 336
 Autoscaling, 329
 Availability, 74, 247, 271, 524
 - cloud services, 9, 260
 Availability zone (AZ), 23
 Average memory access time (AMAT), 51
 AWS, 19, 24
 - networking, 23
 - security, 258, 289, 290
- B**
- Back-end bound, 103
 Back-propagation, 504
 Bad speculation, 103
 Balloon driver, 150
 Balloon process, 104
 Balls-and-bins model, 400
 Bandwidth, 42, 88, 204, 293
 - full bisection, 189
 - memory, 48, 60
 Bare metal hypervisor, 139
 Barrier synchronization, 352, 354, 355
 Base images, 129
 Basic Core Equivalent (BCE), 70
 Basic Paxos, 398
 Benchmarks, 47
 Best-effort delivery, 178
 Biased strategy, 53
 Bid, 205
 Bigquery, 26
 BigTable, 243, 255, 456
 Binary translation, 140
 Bipartite graph, 363
 Bisection width, 189
 Bit-by-bit Round-robin (BR), 204
 Bit-level parallelism, 45
 Black Widow topology, 193, 212
 Bloch sphere, 510
 Block, 61
 - cache, 50, 228
 - data, 50
 Blockchain, 89
 Bootstrap method, 492
 Bootstrap Performance Diagnostic (BPD), 493
 Bootstrap substitution principle, 492
 Border Gateway Protocol (BGP), 26
 Borg, 107
 Borglets, 108
 BorgMaster, 108–110
 Borrowed virtual time (BVT), 150, 311, 347
 Borrowing, 207
 Bound
 - back-end, 103
 - front-end, 103
 Branch condition speculation, 53
 Branch prediction, 53
 Branch target speculation, 53
 Breath first search (BFS), 449
 Bridge, 154
 Broker, 14
 Browny cores, 99
 Bucket, 21, 437, 460, 582
 Buffer space, 204
 Bulk synchronous parallel model, 360
 Bundles, 322
 Bundling, 278
 Business intelligence, 456
 Byte-range tokens, 230
 Byzantine attack, 530
- C**
- Cache, 49
 - banks, 51
 - block, 50, 228
 - capacity, 49, 50
 - data, 251
 - direct map, 50
 - file, 227, 228, 440
 - fully associative, 50
 - hits, 50, 51
 - implementation, 50
 - isolation, 141
 - L1D, 49
 - line, 50, 55
 - miss, 50, 136
 - miss rate, 51
 - n-way set associative, 50
 - on-chip L2, 49
 - row, 50
 - set-associative, 50
 Cache Allocation Technology (CAT), 118
 Cache Array Routing Protocol (CARP), 209
 Cache Monitoring Technology (CMT), 118
 Calculus-based trust, 267
 Callback, 106
 Callstack, 100
 CAP theorem, 88

- Carrier, 14
- Causal message delivery, 378
- Cause–effect relationships, 375
- CDN, 209
 - providers, 209
- Ceil rate (CR), 207
- Cell, 98, 108
 - storage, 220, 395, 396
- Central Limit Theorem (CLT), 495
- Central processing unit (CPU), 358
- Centralized control, 293
- Certificate Authority (CA), 32
- Chain, 404, 576
- Chaining, 59
- Chameleon, 38
- Checkpoint file, 110
- Checkpoint–restart procedures, 350
- Child images, 129
- Chroot, 97, 128, 164
- Chubby files, 236, 243
- Chubby locks, 236
- Chunk, 231
 - column, 439
 - handle, 231
 - servers, 232, 233
- Circular wait, 394
- CISC architecture, 43, 48
- Class of service (COS), 119
- Class-Based Queueing (CBQ), 206
- Classification engine, 115
- Classifier, 206
- Cleanup phase, 430
- Client
 - library, 461
 - thin, 11
- Client–server paradigm, 41, 75–77, 225, 413
- Clock
 - condition, 380
 - processor, 44
- Clos networks, 193
- Cloud
 - application development, 412
 - Bigtable, 26
 - CDN, 26
 - community, 4, 344
 - computer, 4
 - computing
 - delivery models, 14
 - infrastructure, 13, 63, 95, 246, 357, 412, 503
 - controller, 167, 335
 - controls matrix, 263
 - data encryption, 268
 - database, 222
 - DBaaS, 16
 - DNS, 26
 - elasticity, 22, 293, 325, 347, 411, 503
 - functionality, 502
 - Google, 25
 - hardware, 98
 - infrastructure, 455
 - hybrid, 4
 - IBM, 30
 - Identity and Access Management (IAM), 27
 - interconnection networks, 191, 503
 - interoperability, 13, 31, 32, 294, 344
 - KMS, 27
 - Machine Learning, 27, 506
 - Natural Language API, 27
 - oracle, 15
 - private, 4, 5, 167, 168
 - public, 4
 - scheduling subject, 294, 315
 - SDK, 27
 - security, 15, 30, 36, 258, 259, 263, 286, 289–291, 452
 - risks, 259
 - scanner, 27
 - service availability, 9, 260
 - source repositories, 27
 - speech API, 27
 - SQL, 26
 - storage, 26, 216, 255, 257, 289
 - diversity, 30
 - virtual, 316
 - vision API, 27
- Cloud Functions (CF), 26
- Cloud Load Balancing, 26
- Cloud Security Alliance (CSA), 261, 290
- Cloud Service Providers (CSP), 1, 32, 96, 135
- Cloud Virtual Network (CVN), 26
- CloudHSM, 290
- CloudLab, 38
- Cloudlets, 480
- CloudStore, 233, 431
- CloudTrail, 129, 289
- CloudWatch, 21, 22
- Cluster
 - controllers, 167
 - management, 107
 - software, 101
 - system, 110
 - scalable, 133
- Co-scheduling, 354
- Coarse control, 335
- Coarse-grained
 - data-parallel applications, 104
 - locks, 234

- parallelism, 354
- transformations, 121
- Cognitive Radio Networks (CRN), 529
- Cognito, 290
- Colored Petri Nets (CPN), 369
- Colossus File System (CFS), 461
- Column chunk, 439
- Column-family databases, 240
- Combinatorial auctions, 322
- Commission failure, 398
- Commit point, 395
- Committer, 459
- Common data bus (CDB), 52
- Common Object Request Broker Architecture (CORBA), 414
- Communicating sequential processes (CSP), 359, 408
- Communication, 353
 - algorithmic, 354
 - bandwidth, 118, 200, 362
 - channel, 369, 388
 - complexity, 354
 - events, 370, 371, 376, 377, 382
 - infrastructure, 175, 197, 200, 209
 - latency, 24, 87, 99, 353
 - location transparent, 191
 - protocols, 80, 84, 101, 117, 118, 349, 354, 374, 388, 525
 - quantum, 509
- Community cloud, 4
- Complex Instruction Set Computer (CISC), 45
- Compliance Level Agreements (CLA), 34
- Compliant Cloud Computing (C3), 34
- Composability bounds, 86
- Composite services, 450
- Composite task, 415
- Computational models, 358
- Compute Engine (CE), 26
- Computer
 - architecture, 42, 140
 - architecture taxonomy, 46
- Computing
 - grid, 85
 - network-centric, 10
 - utility, 1
 - volunteer, 449
- Concurrency, 349, 350, 353, 358, 361, 387, 388, 392
 - control, 228, 392, 460, 461
 - transparency, 73
- Concurrent events, 376
- Concurrent instructions, 46
- Confidentiality, 271
- Conflict fraction, 112
- Conflicting requirements, 502
- Confusion, 364
 - asymmetric, 364
 - symmetric, 364
- Congestion avoidance (CA), 184
- Congestion control, 180, 183, 185, 211, 344, 375
- Conjugate gradient (CG), 444
- Connection-oriented protocol, 183
- Connectionless protocols, 183
- Consensus service, 397
- Consistency
 - external time, 222
- Consistent cut, 382
- Consistent message delivery, 379
- Consumer service, 14
- Container, 129
- Container Engine (CntE), 26
- Container Registry, 26
- Containers, 95, 97, 128, 130, 434
 - Docker, 26, 128, 129
- Content, 10
 - network-centric, 10
- Content Delivery Network (CDN), 28, 207, 208
- Content service network (CSN), 211
- Content Store, 186
- Content-centric distribution, 211
- Content-Centric Networks (CCNs), 185, 210
- Context switch allowance, 311
- Control
 - coarse, 335
 - concurrency, 228, 392, 460, 461
 - congestion, 375
 - data, 30, 142
 - error, 374
 - fine, 335
 - flow, 43, 356, 374
 - granularity, 335
 - hazard, 46
 - hypervisor, 104
 - integral, 335
 - optimal, 330, 332
 - planes, 10, 161
 - system, 333, 334, 463
 - task, 331
 - theory, 296, 328–330, 333, 335, 346
 - third-party, 261
- Control Data Corporation (CDC), 72
- Control Unit (CU), 44
- Controller, 329, 333, 347, 458
 - application, 335
 - cloud, 167, 335
 - cluster, 167
 - fabric, 29
 - replication, 131
- Convolutional Neural Networks (CNN), 506
- Cooperative spectrum-sensing strategy, 530

- Coordinator, 421
 - Core
 - browny, 99
 - layer, 200
 - network, 176, 209
 - wimpy, 99
 - Correlating predictors, 54
 - Cost
 - memory, 48
 - network, 194
 - Costate equation, 331
 - CPU
 - optimization, 51
 - registers, 49
 - virtualization, 163
 - Critical section, 392, 396, 405, 407
 - Critical vulnerabilities, 54
 - Cron, 107
 - Crunch pipeline, 407
 - CSPs, 1, 15, 32, 96, 135, 335, 344, 345
 - Cumulative distribution function (CDF), 484
 - Current Privilege Level (CPL), 163
 - Current Program Status Register (CPSR), 56, 57
 - Customer, 181
 - Cut, 381
 - consistent, 382
 - frontier, 381
 - Cyberwarfare, 257
 - Cycle per Instruction (CPI), 44
- D**
- D-Wave, 518
 - Daemon, 105, 129
 - Data
 - application, 2, 6
 - block, 50
 - cache, 251
 - control, 30, 142
 - flow, 44, 180, 204, 233, 356, 357, 369, 409
 - architecture, 43, 92
 - generality, 75
 - hazards, 46
 - leakage, 261
 - link layer, 177
 - loss, 261
 - mashups, 450
 - memory flow, 52
 - networks, 210
 - sharing abstraction, 121
 - storage layer, 16, 498
 - streaming, 466, 469, 470
 - applications, 454, 466
 - pipeline, 483
 - services, 483
 - streams, 46
 - tier, 4
 - training, 504
 - warehouse, 456
 - Data Access Units (DAU), 464
 - Data center networks (DCN), 200
 - Data Description Language (DDL), 438
 - Data Manipulation Language (DML), 438
 - Data Mining (DM), 501
 - Data Oriented Network Architecture (DONA), 185
 - Data-aware scheduling, 303
 - Data-level parallelism, 67
 - Data-parallel applications, 99
 - Data-shipping, 230
 - Database, 438
 - cloud, 222
 - column-family, 240
 - document, 240
 - graph, 241
 - layer, 16
 - NoSQL, 222, 239–241, 243, 245, 247, 255, 269, 454
 - provenance, 216, 252
 - schema, 454
 - systems, 247, 269
 - management, 222
 - transactions, 255, 422
 - Database as a Service (DBaaS), 16
 - Datagram, 176
 - Datalog conjunctive query, 253
 - Dataspaces, 28, 499
 - De-perimeterization, 9
 - Deadline
 - absolute, 316
 - relative, 315
 - Deadlocks, 350, 367, 392
 - Decentralized consensus, 90
 - Decision models, 33
 - Decoherence, 513
 - Deep neural networks, 63
 - Deep Neural Networks (DNN), 506
 - Defense Information Systems Agency (DISA), 86
 - Deferred evaluation, 406
 - Degree, 189
 - Delay scheduling, 298, 300
 - Delay slots, 52
 - Delayer, 206
 - Delivery
 - best-effort, 178
 - consistent message, 379
 - First-In-First-Out (FIFO), 378
 - rule, 378
 - Denali hypervisor, 138

- Denial-of-service (DOS), 282, 575
- Deutsch–Jozsa problem, 515
- Device, 61
- DHCP server, 179
- Direct map cache, 50
- Directed acyclic graph (DAG), 125
- Disk locality, 251
- Disk maps, 230
- Disruptive computing paradigm, 2
- Disruptive technology, 39
- Distributed computing fallacies, 88
- Distributed Data Interface (DDI), 443
- Distributed sort, 426
- Distributed system, 73
- Distribution
 - content-centric, 211
 - heavy-tail, 484
 - network, 185
- DNS
 - cloud, 26
 - name, 19
- DNS-based routing, 210
- Docker containers, 26, 128, 129
- Dockerfile, 129
- Domain, 149
 - driver, 154
 - guest, 154
- Domain Specific Architectures (DSA), 62, 507
- DomU, 282
- Double data rate (DDR), 195
- DRAM, 49, 65, 216
 - bandwidth, 117, 118, 120
 - latency, 99
- Driver, 438
 - domain, 154
 - split, 151
- Dynamic
 - application scaling, 328
 - binary translation, 82
 - data-driven applications, 463
 - frequency scaling, 117
 - thresholds, 334
 - trust, 267
 - voltage scaling, 117
 - workflows, 421
- Dynamic Host Configuration Protocol (DHCP), 179, 212
- Dynamic Random Access Memory (DRAM), 216
- Dynamic voltage and frequency scaling (DVFS), 117
- DynamoDB, 22, 247
- E**
- Earliest Deadline First (EDF), 298
- Earliest Virtual Time (EVT), 311
- Ecological impact, 7
- ECS, 129
- Edge
 - algorithm, 126
 - network, 176
- Effective Machine Utilization (EMU), 133
- Effective virtual time, 311
- Elastic Beanstalk, 22
- Elastic Block Store (EBS), 21
- Elastic Compute Cloud (EC2), 19
- Elastic Container Service (ECS), 147
- Elastic Fabric Adapter (EFA), 61
- Elastic IP address, 19, 20
- Elastic Kubernetes Service (EKS), 147
- Elastic Load Balancer (ELB), 10, 22
- Elastic load balancing, 20
- Elastic MapReduce (EMR), 22
- ElastiCache, 22
- Elasticity, 325
- Embarrassingly parallel application, 68, 353
- Embarrassingly parallel problems, 67
- Emergence, 342
- Emulation, 136
 - processor, 137
 - system, 145
 - user-mode, 145
- Enactment engine, 416
- Enactment model, 416
- Energy efficiency, 325, 326
- Energy saving, 327
- Energy use, 7
- Energy-proportional systems, 326
- Enforced modularity, 75
- Enhanced data rated (EDR), 195
- Enterprise flash drive (EFD), 218
- Enterprise Mobile Management (EMM), 287
- Entity groups, 245, 246
- Equal Partitioning Rule (EPR), 319
- Error, 122
 - control, 374
 - sampling, 495
- Estimator, 206
- Eucalyptus, 167
- Event, 369, 376, 470
 - time, 471
- Eventual consistency, 249
- Exception bitmap, 144
- Exception Handler (EH), 388
- Excess vector, 323
- Execute instruction, 52
- Executer functions, 106
- Execution time
 - parallel, 68

- sequential, 68
- Executor, 106, 428
- Extended Nets, 367
- Extensibility, 414
- Extensible hashing techniques, 229
- External Interface, 438

- F**
- Fabric controller, 29
- Facts, 253
- Failure transparency, 73
- Fair Queuing (FQ), 204
- Fair share timeout, 302
- Fair-share scheduler, 127
- Fast Fourier Transform (FFT), 511
- Fast retransmit (FR), 184
- Fat-tree, 192, 193
- Fate sharing, 77
- FauxMaster, 110
- Feasibility, 109
- Feasible state, 332
- Federal Information Processing Standard (FIPS), 261
- Federal Information Security Management Act (FISMA), 261
- Fiber-to-the-home (FTTH), 182
- Fibre Channel (FC), 197
- Field Programmable Gate Arrays (FPGA), 67
- File
 - cache, 227, 228, 440
 - checkpoint, 110
 - Chubby, 236, 243
 - local, 225
 - lock, 236
 - remote, 225, 227
- Filter, 122
- Fine control, 335
- Fine-grained cluster resource sharing, 106
- Fine-grained locks, 234
- Fine-grained parallelism, 66, 354, 355
- Finite state machine (FSM), 53
- Firing sequence, 367
- First-In-First-Out (FIFO) delivery, 378
- Fixed parallel time per process, 69
- Fixed problem size, 69
- Flash crowd, 209
- Flash memory, 217
- Flattened butterfly network, 193
- Floating-point performance models, 60
- Flow
 - control, 183, 184, 189, 197, 296, 356, 374
 - data, 44, 180, 204, 233, 356, 357, 369, 409
 - memory, 52
 - register, 52
 - instruction, 52
- relations, 366
- Flow-control mechanism, 179
- Flowchart, 416
- FlumeJava, 406, 407, 428
 - library, 406
- Fork routing task, 416
- Forwarding Information Base, 187
- Forwarding plane, 178
- Front-end bound, 103
- Frontier of the cut, 381
- Full bisection bandwidth, 189
- Full virtualization, 135, 140
- Fully associative cache, 50
- Fully Homomorphic Encryption (FHE), 268
- Function
 - aggregation, 458
 - objective, 322
 - utility, 338
 - watch, 151

- G**
- Gather operation, 59
- General Parallel File System (GPFS), 229, 255
- Generalized Processor Sharing (GPS), 339
- Generic contraction algorithm, 556
- Generic TLC algorithm, 554
- Geo-replication, 457
- Geometric engine, 67
- GFS files, 231
- Global predicate evaluation problem (GPE), 374
- Global real-time clock, 379
- Gmail, 27
- Google
 - Base, 28
 - Calendar, 26, 27, 34, 37
 - Cloud, 25
 - infrastructure, 25, 98
 - Co-op, 28
 - Docs, 27
 - Drive, 26, 28, 39
 - Groups, 28
 - Maps, 28, 37
- Google App Engine, 26
- Google Container Engine (GKE), 129
- Google file system (GFS), 216, 256
- Gradient, 304
 - descent, 504
- Grand architectural complications, 48
- Graph
 - bipartite, 363
 - databases, 241
 - marked, 365, 368
- Graphics Processing Unit (GPU), 2, 60, 67

- GrepTheWeb application, 428, 451
 Ground truth, 493
 Guest
 - domain, 154
 - hypervisor, 142, 156, 158, 159
 - state, 142, 143
- H**
- Hadoop, 115, 431, 435
 - Apache, 105, 320, 428, 431, 432
 - master, 298
 - scheduler, 298, 302
 - system, 431
 - tasks, 299, 320
 - workloads, 115
- Hadoop Distributed File System (HDFS), 105
 Hadoop fair scheduler (HFS), 302
 Hard disk, 49
 Hard disk drive (HDD), 216, 217
 Hardware
 - cloud, 98
 - costs, 137, 171, 223
 - parallel, 68
 - parallelism, 42
 - support, 136, 141
 - virtual, 97, 137, 169
 - virtualization, 146, 171
- Hardware description language (HDL), 67
 Hashing
 - perfect random, 404
- Hazard
 - control, 46
 - data, 46
 - pipeline, 46
 - structural, 46
- HDFS, 105, 431
 Head node, 316
 Heavy-tail distributions, 484
 Heracles, 119
 Hidden states, 143
 Hiding the latency, 42, 43, 461
 Hierarchical coordination scheme, 421
 Hierarchical Token Buckets (HTB), 207
 Hierarchy, 79
 - of networks, 98
- High Performance Computing (HPC), 432
 History-based prediction, 53
 Hit time, 51
 Hive, 437
 Homomorphic encryption, 269
 Horizontal scaling, 242, 329
 Host, 176
 - hypervisor, 156
 - state, 142, 143
- Hosted, 139
 - independent hypervisors, 171
 - specialized hypervisors, 171
- Hosting
 - KVM, 145
 - Xen, 145
- Hostname, 19
 Hourglass architecture, 178
 HTTP redirection, 210
 HTTP server, 77
 HTTP-tunneling, 79
 Hub, 129
 Hybrid, 139
 - cloud, 4
- Hyper-threading, 47, 67
 Hypercalls, 150
 Hypervisor
 - bare metal, 139
 - Denali, 138
 - guest, 142, 156, 158, 159
 - host, 156
 - hosted independent, 171
 - hosted specialized, 171
 - native, 170
 - rogue, 169
 - Xen, 19, 138, 155, 164, 172, 281
- I**
- IBM clouds, 30
 Image, 129
 - base, 129
 - child, 129
 - official, 129
 - user, 129
- Impala, 438
 In-memory cluster computing, 120
 Incidence Matrix, 367
 Incommensurate scaling, 343
 Independence, 414
 Index vector, 59
 InfiniBand, 194, 195
 Information, 10
 - hiding, 74
- Information Management System (IMS), 254
 Information Processing Technology Office (IPTO), 93
 Infrastructure
 - cloud computing, 13, 63, 95, 246, 357, 412, 503
 - communication, 175, 197, 200, 209
 - Google Cloud, 25, 98
 - network, 82, 181
 - networking, 94, 191
 - virtualization, 136, 145

- Infrastructure as a Service (IaaS), 16
 - Infrastructure as Code (IaC), 18
 - Inhibitor arc, 367
 - Initial events, 359
 - Insecure APIs, 261
 - Inspector (IS), 289
 - Instance, 547
 - Instruction
 - concurrent, 46
 - flow, 52
 - issue, 52
 - privileged, 81, 162
 - privileged-sensitive, 163
 - sensitive, 81, 173
 - Instruction Access Units (IAU), 464
 - Instruction Level Parallelism (ILP), 101
 - Instruction Register (IR), 43, 163
 - Instruction Set Architecture (ISA), 43, 82
 - Instruction-level parallelism (ILP), 45
 - Instructions per Clock Cycle (IPC), 44
 - Integral control, 335
 - Integrated circuits (IC), 92
 - Intercloud, 32, 33, 345
 - exchange, 32
 - root, 32
 - nodes, 32
 - security, 291
 - Interconnection network, 47, 70, 85, 176, 189, 192, 194, 222, 404, 432, 464, 482
 - for computer clouds, 189
 - Interdependence, 267
 - Interface, 81
 - external, 438
 - network, 24, 178, 179, 189, 478, 480
 - Interface Definition Language (IDL), 414
 - Interface Message Processors (IMP), 93
 - Internal clock, 375
 - International Data Corporation (IDC), 7
 - Internet, 93, 175–178
 - transport protocols, 179
 - Internet Assigned Numbers Authority (IANA), 182
 - Internet Cache Protocol (ICP), 209
 - Internet exchange point (IXP), 181
 - Internet protocol (IP), 8
 - Internet Research Task Force (IRTF), 186
 - Internet Service Providers (ISP), 181
 - Interrupt virtualization, 143
 - Interval timer, 375
 - Intrusion detection, 268
 - systems, 575
 - Invariant behavior, 74
 - IP address, 178
 - elastic, 19, 20
 - private, 19
 - public, 19
 - IP anycasting, 210
 - IPsec, 183
 - Iso-latency policy, 116
 - Isolation
 - cache, 141
 - modules, 106
 - performance, 137
 - security, 37, 135, 137, 412
 - software fault, 450
 - Itanium
 - architecture, 142, 161, 173
 - paravirtualization, 136
 - processor, 162, 163
 - Iterator method, 121
 - IVA register, 163
- J**
- Java, 405
 - threads, 405
 - JAVA Message Service (JMS), 450
 - Java Virtual Machine (JVM), 273, 405
 - Job Manager (JM), 105
 - instance, 447
 - Job registry, 447
 - Job tracker, 298, 431, 432
 - Join, 121, 122
 - point, 465
 - routing task, 416
 - Journal storage, 220, 395
 - Just-in-time infrastructure, 411
- K**
- Kernel-based Virtual Machine (KVM), 145
 - Key Management Interoperability Protocol (KMIP), 33
 - Key Management Service (KMS), 268, 290
 - Key-based routing (KBR), 84
 - Key-value, 240, 241
 - Killer application, 37
 - Kimberlay, 481
 - Knowledge base (KB), 267
 - Kubernetes, 130–132
 - KVM hosting, 145
 - KVM virtualization, 148
- L**
- L1D cache, 49
 - Labels, 131
 - Lagrange multiplier method, 330
 - Lambda service, 23
 - Language INtegrated Query (LINQ), 105
 - Large Hadron Collider (LHC), 85

- Large-scale systems, 84
 - Last level cache (LLC), 118
 - Late binding, 303
 - Latency, 42, 293, 479, 484
 - communication, 24, 87, 99, 353
 - constraints, 474
 - DRAM, 99
 - hiding, 42, 43, 461
 - memory, 48, 67, 101, 103, 294, 443
 - Latency Critical (LC), 116
 - Latency-critical workloads, 116, 117
 - Layer, 81
 - access, 200
 - aggregate, 200
 - application, 16, 80
 - core, 200
 - data link, 177
 - data storage, 16, 498
 - database, 16
 - lineage, 127
 - network, 80, 176, 177
 - persistence, 127
 - processor abstraction, 163
 - transport, 80, 176
 - user interface, 16
 - Layering, 79, 176
 - Leader, 398, 423
 - participant, 460
 - Learner, 398
 - Learning rate, 114
 - Least Recently Used (LRU), 228
 - LHC Computing Grid (LCG), 85
 - Lineage, 121
 - layer, 127
 - Links, 122
 - Live media, 209
 - Live migration, 104, 146
 - Livelock condition, 394
 - Livelock freedom, 407
 - Liveness, 367, 416
 - Load balancing, 13
 - Load-store architecture, 43
 - Local
 - file, 225
 - memory, 251, 303, 462
 - relevance, 211
 - snapshot, 236
 - Locality, 48, 299, 455
 - disk, 251
 - rack, 299
 - servers, 299
 - spatial, 49, 356
 - temporal, 49, 356
 - Location transparency, 73, 227
 - Location transparent communication, 191
 - Lock, 392, 396
 - advisory, 234
 - Chubby, 236
 - coarse-grained, 234
 - file, 236
 - fine-grained, 234
 - granularity, 230
 - mandatory, 234
 - table, 460
 - tokens, 230
 - Locking, 234
 - Log data centers, 483
 - Logical clock (LC), 376
 - Logical organization, 223, 224
 - Logical Plan (LP), 498
 - Long Range Computer Study Group (LRCSG), 92
 - Long Short-Term Memory (LSTM), 506
 - Loopback file system (LOFS), 278
- M**
- Machine Intelligence (MI), 501
 - Machine Learning (ML), 11, 25, 63, 296, 501, 503, 506
 - Machine Perception (MP), 501
 - Machine-based automation, 90
 - Main memory, 49
 - Main Stack Pointer (MSP), 56
 - Malicious insiders, 261
 - Malicious software, 169
 - Malware, 169
 - Management
 - QoS-aware cluster, 113
 - resource, 127
 - shared state cluster, 111
 - Mandatory
 - locks, 234
 - security, 272
 - Map, 121, 122
 - Mapping a computation, 329
 - MapReduce, 105, 425, 426, 436, 449, 451
 - application, 113, 123, 133, 320, 433, 435, 452
 - scheduling, 320
 - framework, 433
 - Marked graph, 365, 368
 - Marking, 363
 - Markov decision processes, 488
 - Master, 232, 235, 236, 240, 425
 - Hadoop, 298
 - lease, 235
 - Master-slave replication, 240
 - Max-min fairness criterion, 297
 - Mechanism, 294

- flow-control, 179
 - parallel recovery, 467
 - Media Access Control (MAC), 80
 - Megastore, 245, 255, 438
 - Meltdown vulnerability, 54
 - Memcaching, 242
 - Memory
 - access, 54, 61, 138, 360, 361, 455
 - bandwidth, 48, 60
 - banks, 59
 - cost, 48
 - latency, 48, 67, 101, 103, 294, 443
 - local, 251, 303, 462
 - virtual, 51, 136, 155, 156, 171, 228, 250, 388, 409
 - virtualization, 104, 164, 172
 - Memory Management Unit (MMU), 57, 104, 161
 - Mesa, 106, 457, 458
 - Message, 356, 425
 - Message delivery
 - causal, 378
 - rules, 377
 - Message Passing Interface (MPI), 388
 - Metanode, 230
 - Metaphysical addressing, 164
 - Method
 - bootstrap, 492
 - iterator, 121
 - Lagrange multiplier, 330
 - Microservices, 18
 - Microsoft Azur Container Service, 130
 - Microsoft Windows Azure, 28
 - Middleware, 73, 85
 - Migration transparency, 74
 - MIMD, 47
 - Minimal witness basis, 253
 - Minimum charging units, 311
 - Minimum share timeout, 302
 - Minions, 130, 131
 - Miss penalty, 51
 - Mitigating cloud vulnerabilities, 287
 - Mobile Application Management (MAM), 287
 - Mobile Device Management (MDM), 287
 - Mode
 - overlay, 530
 - sensitive, 140
 - Model
 - arbitrarily divisible load sharing, 412
 - assumptions, 534
 - balls-and-bins, 400
 - bulk synchronous parallel, 360
 - computational, 358
 - decision, 33
 - enactment, 416
 - floating-point performance, 60
 - for multicore computing, 361
 - navigation, 254
 - NoSQL distribution, 240
 - roofline, 60
 - storage, 220
 - three-tier, 4
 - weak coordination, 421
 - Modularity, 74
 - Modularization, 176
 - Modularly divisible application, 329
 - Monitor, 334
 - Monitoring phase, 430
 - Multi-Layer Perceptrons (MLP), 506
 - Multicasting, 183
 - Multicloud, 39
 - Multicore processor speedup, 70
 - Multilevel nested virtualization, 156
 - Multiplexing, 136
 - Multistage interconnection networks, 193
 - Multitenancy, 259
 - Multithreaded SIMD processors, 61, 65
 - Mutual exclusion, 407
 - Myrinet, 194, 196
- N**
- N-way set associative cache, 50
 - Name node, 431, 432
 - Name Server (NS), 105
 - Named Data Networks (NDN), 186
 - National Science Foundation (NSF), 38
 - Native hypervisors, 170
 - Natural language processing (NLP), 507
 - Natural language understanding (NLU), 507
 - Navigational model, 254
 - Negotiation
 - object, 33
 - protocols, 33
 - Nested virtualization, 136, 142, 156, 158, 159, 172
 - multilevel, 156
 - single-level, 156
 - Network
 - adapter, 154
 - architecture, 154, 176
 - Clos, 193
 - congestion, 183, 204
 - content service, 211
 - content-centric, 185, 210
 - convolutional neural, 506
 - core, 176, 209
 - cost, 194
 - deep neural, 506
 - diameter, 189

- distribution, 185
- edge, 176
- fat-tree, 193
- flattened butterfly, 193
- infrastructure, 82, 181
- interconnection, 47, 70, 85, 176, 189
- Interconnection, 189
- interconnection, 189, 192, 194, 222, 404, 432, 464, 482
 - cloud, 191, 503
- interface, 24, 178, 179, 189, 478, 480
- layer, 80, 176, 177
- model, 238, 344
- multistage interconnection, 193
- overlay, 84, 188, 211
- packet-switched, 176
- resource management algorithms, 204
- software-defined, 187
- storage area, 197
- tensor, 553
- Tier 1, 181
- Tier 2, 181
- Tier 3, 181
- topology, 189, 231
- vehicular ad hoc, 211
- virtual, 84, 212, 480, 481
- virtualization, 154, 156
- Network Address Translation (NAT), 19, 182
- Network Element Control Protocol (NECP), 209
- Network File System (NFS), 76, 88, 216, 221, 223
- Network Interface Card (NIC), 153
- Network-centric computing, 10
- Network-centric content, 10
- Networking
 - AWS, 23
 - infrastructure, 94, 191
- Networks-on-a-chip (NoC), 65
- Neural networks (NN), 504
- Neutrality, 414
- Next Generation Enterprise Network (NGEN), 86
- Node, 189, 320
 - head, 316
 - name, 431, 432
 - worker, 316
- Non Secure World (NSW), 160
- Nonpreemption, 394
- Nonrepudiation, 32
- NoSQL databases, 222, 239–241, 243, 245, 247, 255, 269, 454
- NoSQL distribution models, 240
- O**
- Oak Ridge National Laboratory (ORNL), 72
- Object code, 82
- Object of negotiation, 33
- Object Request Broker (ORB), 414
- Object-relational impedance mismatch, 239
- Objective function, 322
- Official images, 129
- Omission failure, 398
- On-chip L1I, 49
- On-chip L2 cache, 49
- ONF architecture, 188
- Ontology, 32
- Opcode, 59
- Open Container Initiative (OCI), 130
- Open Network Foundation (ONF), 188
- Open Shortest Path First (OSPF), 187
- Open-Nebula, 168
- OpenStack, 168
- Operating system (OS), 57, 96, 135
 - security, 272
- Operation
 - gather, 59
 - scatter, 59
- Optimal control, 330, 332, 346, 490
- Optimal Partitioning Rule (OPR), 316
- Optimal resource management, 329
- Oracle, 515
- Order Preserving Encryption (OPE), 270
- Organization
 - kubernetes, 130
 - logical, 223, 224
 - physical, 223
 - scale-free, 344
- Over-provisioning, 325
- Overclocking, 118
- Overlay, 209
 - links, 84
 - networks, 84, 188, 211
- Overlimit, 207
- Overprovisioning, 6
- Oversubscription, 98, 191
- P**
- P2P replication, 240
- PaaS, 15, 16, 26, 34, 450
- Packet filtering, 575
- Packet-switched networks, 176
- Packets, 176, 425
- Pages, 51
- Paradigm
 - client-server, 41, 75–77, 225, 413
 - disruptive computing, 2
- Parallel
 - collection, 406
 - execution time, 68
 - file systems, 216, 222, 228

- hardware, 68
- recovery mechanism, 467
- slackness, 360
- Parallel File Systems (PFS), 221
- Parallel random-access machine (PRAM), 358
- Parallelism, 45
 - bit-level, 45
 - coarse-grained, 354
 - data-level, 67
 - fine-grained, 66, 354, 355
 - instruction-level, 45
 - task-level, 67
 - Thread-level, 67
- Paramount complication, 48
- Paravirtualization, 135, 140, 141, 148, 161, 164, 171, 274
 - Itanium, 136
 - strategy, 153
 - Xen, 149
- Partial Differential Equations (PDE), 66, 369
- Participant leader, 460
- Pass-through operators, 498
- Passphrase, 280
- Path, 224
 - trusted, 272
- Pattern, 151
 - workflow, 419
- Paxos, 397
 - algorithm, 236, 237
 - basic, 398
 - group, 460
- Peer-to-peer systems, 82
- Peering, 181
- Pending Interest Table (PIT), 187
- Perceptron, 504
 - rule, 505
- Perfect random hashing, 404
- Performance
 - metrics, 47, 113, 210, 338
 - transparency, 74
 - unpredictability, 137
 - VM, 153
 - WSC, 100
- Periodic task, 315
- Persistence layer, 127
- Persistent backlog, 207
- Persistent trust, 267
- Petri nets (PN), 363
- Phase
 - cleanup, 430
 - monitoring, 430
 - shutdown, 430
 - transition, 343
- Photon, 474
- Physical organization, 223
- Physical Plan (PP), 498
- Picasa, 27
- Pig, 436
- Pinhole, 575
- Pipeline
 - hazards, 46
 - scheduling, 46
 - stall, 46
- Pipelining, 45
- Plane
 - forwarding, 178
 - routing, 179
- Planning, 416
- Platform as a Service (PaaS), 15
- Pod, 201
- Point-of-presence (POP), 181
- Policy, 268, 294
- Port, 179, 180
- Precision, 375
- Precursors, 223
- Predecessor task, 415
- Prediction
 - branch, 53
 - history-based, 53
- Presentation tier, 4
- Preserve exception behavior, 48
- Preserve instruction flow, 48
- Preserve security, 48
- Pricing, 322
- Primitive processes, 359
- Primitive task, 415
- Priority bands, 109
- Priority inheritance, 127
- Priority inversion, 394
- Priority-based scheduling, 127
- Privacy
 - impact, 291
 - risks, 266, 278, 280
- Privacy Impact Assessment (PIA), 264, 266
- Private
 - cloud, 4, 5, 167, 168
 - IP address, 19
- Privilege leaking, 162
- Privilege rings, 162
- Privileged instructions, 81, 162
- Privileged-sensitive instructions, 163
- Probability density function (PDF), 484
- Problem
 - Deutsch–Jozsa, 515
 - embarrassingly parallel, 67
- Process, 82, 359, 369
 - balloon, 104

- description, 416
 Markov decision, 488
 primitive, 359
 VM, 138
- Process Stack Pointer (PSP), 56
- Processor
 abstraction layer, 163
 asymmetric core, 70
 bandwidth, 43, 307, 356
 clock, 44
 emulation, 137
 Itanium, 162, 163
 management, 81
 multithreaded SIMD, 61, 65
 resources, 117
 sharing, 96
 superscalar, 45, 47
 symmetric core, 70
 utilization, 96, 132, 336, 353, 361
 virtualization, 80, 96, 135, 137, 140, 172, 388, 409
- Producer-consumer synchronization, 407
- Program counter (PC), 43
- Proportional thresholding, 335
- Proposal, 398, 425
- Proposer, 398
- Protocol, 176, 209, 371
 connection-oriented, 183
 connectionless, 183
 negotiation, 33
- Provider
 CDN, 209
 service, 14
- Public cloud, 4
- Public IP address, 19
- Q**
- QEMU, 144
- QoS, 113
 crosstalk, 137
- QoS-aware cluster management, 113
- Quad Data Rate (QDR), 195, 445
- Quantum communication, 509
- Quantum Computing (QC), 501, 509
 Playground, 516
- Quantum Information Processing (QIP), 509
- Quantum Inspire, 518
- Quantum parallelism, 510
- Quasar, 114
- Qubit, 509
 control, 511
- Quorum, 398
- Quota system, 109
- R**
- Rack, 98
 locality, 299
- Random-access machine (RAM), 358
- Rank, 121, 122
- Rate Monotonic Algorithms (RMA), 298
- Reachability, 367
- Read after Write (RAW), 46
- Real-time bandwidth, 118
- Recurrent Neural Networks (RNN), 506
- Recursive construction, 75
- Red Hat Virtualization (RHV), 170
- Regions, 23
- Register
 CPU, 49
 data flow, 52
 stack engine, 163
 vector length, 59
 vector mask, 59
- Register File (RF), 44
- Registration Authority (RA), 32
- Regularization factor, 114
- Relational algebra, 238
- Relational Database Service (RDS), 10
- Relational trust, 267
- Relative deadline, 315
- Remote File System (RFS), 223, 255
- Remote Procedure Call (RPC), 73, 76, 225
- Replanning, 415
- Replay log, 236
- Replication
 controllers, 131
 master-slave, 240
 P2P, 240
 transparency, 73
- Representational State Transfer (REST), 414
- Reputation, 268
- Request for quotation (RFQ), 93
- Request-parallel applications, 99
- Reservation stations (RS), 52
- Resilient Distributed Dataset (RDD), 121, 466
- Resource
 allocation, 113
 assignment, 113
 isolation, 116
 management, 127, 328
 multiplexing, 2
 processor, 117
 scale out, 115
 scale up, 115
 virtualization, 136
- Resource Description Framework (RDF), 85
- Response time, 77, 78, 165

- Ring, 151
 - aliasing, 143
 - compression, 143, 162
 - deprivileging, 142
 - privilege, 162
 - Risk, 267
 - privacy, 266, 278, 280
 - profile, 261
 - security, 259, 271, 279–281, 287, 291
 - Rogue hypervisor, 169
 - Roofline model, 60
 - Root, 97, 576
 - Rootkit, 169
 - Routing
 - DNS-based, 210
 - plane, 179
 - task, 415
 - Row
 - cache, 50
 - groups, 439
 - Rule, 253, 576
 - Run, 381
- S**
- Safety, 416
 - Same Program Multiple Data (SPMD), 67, 425
 - Sample estimate result, 495
 - Sampling, 121
 - Sampling error, 495
 - distribution, 495
 - Scalability, 86, 239, 240
 - Scalable data center communication architectures, 200
 - Scale-free organization, 344
 - Scaled speedup, 69
 - Scaling
 - horizontal, 242, 329
 - incommensurate, 343
 - transparency, 74
 - vertical, 328
 - Scan consolidation, 498
 - Scatter operation, 59
 - Scheduler
 - Apache capacity, 306
 - fair-share, 127
 - Hadoop, 298, 302
 - Spark, 121, 122, 306
 - thread, 63, 311
 - workload, 111
 - Scheduler virtual time (SVT), 311
 - Scheduling, 316
 - data-aware, 303
 - delay, 298
 - Schema, 454
 - Schema-on-write, 454
 - Scoring, 109
 - Search time, 217
 - Searchable symmetric encryption (SSE), 270
 - Secondary spectrum-data falsification (SSDF), 530
 - Secure arguments, 75
 - Secure Sockets Layer (SSL), 183, 262
 - Secure Virtual Machine (SVM), 142
 - Security, 258
 - assessment, 264
 - audit, 15
 - AWS, 258, 289, 290
 - cloud, 15, 30, 36, 258, 259, 263, 286, 289–291, 452
 - Intercloud, 291
 - isolation, 37, 135, 137, 412
 - mandatory, 272
 - of database services, 270
 - operating system, 272
 - risks, 259, 271, 278–281, 287, 291
 - virtual machine, 273
 - virtualization, 257, 275
 - VM, 257, 273
 - Security as a service (SecaaS), 264
 - Security Extensions (SE), 161
 - Seek time, 217
 - Segments, 179
 - Selector, 206
 - Self-organization, 342, 344
 - Semantic integration, 28
 - Semantic Web, 85
 - Sensitive instructions, 81, 173
 - Sensitive machine instructions, 140
 - Sensors, 329, 333, 334
 - Sequencer, 234
 - Sequential execution time, 68
 - Serialization, 392
 - Server
 - chunk, 232, 233
 - HTTP, 77
 - locality, 299
 - stateless, 413
 - thrift, 438
 - utilization, 127, 210, 227, 259, 304, 325
 - virtual, 18, 19, 488
 - WSC, 99
 - Service, 131
 - composite, 450
 - consensus, 397
 - consumer, 14
 - hijacking, 261
 - Lambda, 23
 - provider, 14
 - Service Level Agreement (SLA), 6, 13, 33, 36, 260

- Service Level Objectives (SLO), 116
- Servless computer service, 23
- Sessions, 470
- Set-associative cache, 50
- Sharding, 240
- Shared cluster state, 111
- Shared memory system (SMS), 464
- Shared security responsibility (SSR), 17, 289
- Shared state cluster management, 111
- Shared technology, 261
- Shared write lock, 230
- Shutdown phase, 430
- Sigmoid neuron (SN), 505
- SIMD, 47
 - architectures, 58–60
- Simple DB, 21
- Simple Object Access Protocol (SOAP), 414
- Simple Queue Service (SQS), 21
- Simple Storage System (S3), 20
- Simple Workflow Service (SWS), 22
- Simultaneous multithreading (SMT), 103
- Single data rate (SDR), 195
- Single-level nested virtualization, 156
- Singular Value Decomposition (SVD), 113
- SISD processors, 46
- Slack, 27, 119, 329
- Slave threads, 465
- Sliding, 470
- Slots, 106, 320
 - sticky, 299
- Socket, 88, 179
- Soft modularity, 75
- Soft-state, 241
- Software, 87
 - architecture, 4
 - malicious, 169
 - virtualization, 136, 170
 - vulnerability, 278
- Software as a Service (SaaS), 15
- Software fault isolation (SFI), 450
- Software-defined Network (SDN), 187
- Software-Defined Wide Area Network (SD-WAN), 188
- Solid-state disks, 217
- Space
 - active, 421
 - tuple, 421
- Space–time diagram, 377, 381
- Spanner, 459, 460
- Spark, 121, 122
 - dependencies, 122
 - driver programs, 122
 - scheduler, 121, 122, 306
 - streaming, 466, 467, 470, 498
- Spatial locality, 49, 356
- Spectre, 54
- Speculation
 - branch condition, 53
 - branch target, 53
- Speedup, 68, 353
- Spintronics, 216
- Split, 427
 - drivers, 151
- Sprite File System (SFS), 216, 227, 255
- SRAM, 49
- Stack pointer, 387
- Start-time fair queuing (SFQ), 307, 347
- State, 369
 - feasible, 332
- State Machine, 368
- State of a case, 416
- Stateless address auto-configuration (SLAAC), 183
- Stateless servers, 413
- Static content, 209
- Sticky slots, 299
- Stochastic Fairness Queuing (SFQ), 206
- Stochastic Gradient Descent (SGD), 114
- Stochastic Petri Nets (SPN), 369
- Storage
 - cell, 220, 395, 396
 - cloud, 26, 216, 255, 257, 289
 - journal, 220, 395
- Storage area network (SAN), 197, 221, 255
- Storage manager, 396
- Storage models, 220
- Straggler, 303
- Strategy
 - biased, 53
 - cooperative spectrum-sensing, 530
- Streaming SIMD extensions (SSE), 59
- Stride, 59
- Strong clock condition, 380
- Structural hazards, 46
- Structured Query Language (SQL), 238
- Stubs, 414
- Successor task, 415
- Supernodes, 84
- Superscalar processor, 45, 47
- Supersteps, 360
- Symmetric confusion, 364
- Symmetric core processor, 70
- Synchronization, 350
 - barrier, 352, 354, 355
 - producer–consumer, 407
- System
 - distributed, 73
 - emulation, 145

- energy-proportional, 326
 - intrusion detection, 575
 - large-scale, 84
 - management
 - cluster, 110
 - resource, 75
 - model, 316
 - peer-to-peer, 82
 - quota, 109
 - VM, 138
- T**
- Table schema, 458
 - Tablet, 459
 - Tachyon, 124, 127
 - Tachyon Master (TM), 125
 - Tag, 50
 - Target qubit, 511
 - Task, 360, 415, 416
 - aperiodic, 316
 - composite, 415
 - periodic, 315
 - predecessor, 415
 - primitive, 415
 - routing, 415
 - fork, 416
 - join, 416
 - successor, 415
 - tracker, 298, 431, 432
 - Task-level parallelism, 67
 - Task-priority register (TPR), 143
 - TCP congestion control, 183
 - Temporal locality, 49, 356
 - Tensor, 62
 - contraction, 552, 553
 - network, 553
 - Tensor Processing Units (TPU), 2, 506, 507
 - Terraform, 18
 - Tez, 434
 - Thin clients, 11, 41
 - Third Party Auditors (TPA), 271
 - Third-party control, 261
 - Thread, 67
 - allocate, 385
 - blocks scheduler, 61
 - groups, 354, 369
 - pool, 405
 - scheduler, 63, 311
 - block, 63
 - slave, 465
 - Thread-level parallelism, 67
 - Three-tier model, 4
 - Three-way handshake, 77
 - Thresholds, 335
 - Thrift Server, 438
 - Thunk, 145
 - Tier
 - application/logic, 4
 - data, 4
 - presentation, 4
 - Tier 1 network, 181
 - Tier 2 network, 181
 - Tier 3 network, 181
 - Time
 - intervals, 375
 - response, 77, 78, 165
 - search, 217
 - seek, 217
 - skew, 470
 - stamp, 375, 379, 468
 - Time-slice, 298
 - Timers, 468
 - Timing, 204
 - Token manager, 230
 - Tokens, 366
 - byte-range, 230
 - lock, 230
 - Total surplus, 322
 - TPUs, 2, 42, 63, 506, 507
 - Tracker
 - job, 298, 431, 432
 - task, 298, 431, 432
 - Traditional threats, 260
 - Training data, 504
 - Transaction manager, 460
 - Transit, 181
 - centers, 23
 - Transition, 363
 - phase, 343
 - Transparency
 - access, 73
 - concurrency, 73
 - failure, 73
 - location, 73, 227
 - migration, 74
 - performance, 74
 - replication, 73
 - scaling, 74
 - Transport layer, 80, 176
 - Transport Layer Security (TLS), 183, 282, 287
 - Tread scheduler, 61
 - Triggers, 22, 471
 - Trust, 267, 530
 - calculus-based, 267
 - dynamic, 267
 - indexes, 32

- persistent, 267
- relational, 267
- Trust Zone Secure World (TZSW), 161
- Trusted applications, 272
- Trusted Computer Base (TCB), 38, 281
- Trusted Execution Environment (TEE), 160
- Trusted paths, 272
- Tunneling, 200
- Tuple calculus, 255
- Tuple space, 421
- Two-level resource allocation strategy, 333

- U**
- Undelete files, 278
- Underlimit, 207
- Union, 121
- Universal gates, 505
- Universal Serial Bus (USB), 153
- UNIX File System (UFS), 223, 236
- User
 - challenges, 34
 - images, 129
- User-mode emulation, 145
- Utility computing, 1
- Utility function, 338
- Utility-based, 296

- V**
- Validation Authority (VA), 32
- Values, 457
- Variability, 502
- Vast amounts, 476
- Vector
 - architectures, 59
 - clocks, 249
 - length registers, 59
 - mask registers, 59
- Vehicular
 - clouds, 518, 519
 - networks, 519
- Vehicular ad hoc network (VANET), 211
- Vehicular Virtual Machine Migration (VVMM), 522
- Vendor lock-in, 9, 30
- Vertical scaling, 242, 328
- Very Long Instruction Word (VLIW), 161
- Virtual
 - address, 51
 - cloud, 316
 - hardware, 97, 137, 169
 - memory, 51, 136, 155, 156, 171, 228, 250, 388, 409
 - network, 84, 212, 480, 481
 - server, 19
 - servers, 18, 488
 - time, 311
- Virtual Block Device (VBD), 151
- Virtual circuits (VC), 199
- Virtual Machine Control Structure (VMCS), 142
- Virtual machine (VM), 18, 20, 95, 135, 138, 164, 213
- Virtual Memory Manager (VMM), 388
- Virtual Network Computing (VNC), 481
- Virtual Network Interfaces (VNI), 151
- Virtual Private Cloud (VPC), 22
- Virtual Private Network (VPN), 22, 26
- Virtual-Machine Based Rootkit (VMBR), 169, 170
- Virtualization
 - by aggregation, 95
 - by multiplexing, 95
 - CPU, 163
 - full, 135, 140
 - hardware, 146, 171
 - support, 141
 - interrupt, 143
 - KVM, 148
 - memory, 104, 164, 172
 - nested, 136, 142, 156, 158, 159, 172
 - multilevel, 156
 - single-level, 156
 - network, 154, 156
 - processor, 80, 96, 135, 137, 140, 172, 388, 409
 - resource, 136
 - security, 257, 275
 - software, 136, 170
- VM
 - application, 138
 - import, 20, 547
 - performance, 153
 - process, 138
 - security, 257, 273
 - system, 138
- Volunteer computing, 449
- von Neumann architecture, 91
- Vulnerability
 - critical, 54
 - software, 278

- W**
- Wait time, 302
- Warehouse Scale Computers (WSC), 95, 96, 98
- Watch function, 151
- Wave-particle duality principle, 508
- Weak coordination models, 421
- Web Application Firewall (WAF), 290
- Web browser, 77
- Web Cache Coordination Protocol (WCCP), 209
- Web page, 77

- Web Services Description Language (WSDL), 414
Wide Area Networks (WAN), 73
Wimpy cores, 99
Windowing, 470
Windows Azure, 28
Wireless local area networks (WLAN), 478
Witness tree, 404
Worker node, 316
Worker role, 29
Workflow, 415
 dynamic, 421
 pattern, 419
Working set, 49, 120, 455
Workload
 derivative, 316
 latency-critical, 117
 scheduler, 111
Write after Read (WAR), 46
Write after Write (WAW), 46
Write-ahead, 229
Write-write conflict, 240
WSC
 performance, 100
 servers, 99
- X**
- XDP, 77
Xen, 148
 hosting, 145
 hypervisor, 19, 138, 155, 164, 172, 281
 paravirtualization, 149
XenStore, 151, 282, 284, 285, 291
Xoar, 283
- Y**
- Yarn, 433
- Z**
- Znode stores, 425
Zonemaster, 459